Efficient Implementation Strategies for Block Ciphers on ARMv8

Bachelorarbeit

Bastian Engel

March 23, 2023

Abstract

Lorem ipsum dolor [1] sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Declaration

I hereby declare that ...

Contents

Chapter 1

Introduction

1.1 Notation

1.2 Block ciphers

Securing communication channels between different parties has been a long-term subject of study for cryptographers and engineers which is essential to our modern world to cope with ever-increasing amounts of devices producing and sharing data. The main way to facilitate high-throughput, confidential communications nowadays is through the use of symmetric cryptography in which two parties share a common secret, called a key, which allows them to encrypt, share and subsequently decrypt messages to achieve confidentiality against third parties. Ciphers can be divided into two categories; block ciphers, which always encrypt fixed-sized messages called blocks, and stream ciphers, which continuously provide encryption for an arbitrarily long, constant stream of data.

A block cipher can be defined as a bijection between the input block (the message) and the output block (the ciphertext). For any block cipher with block size n, we denote the key-dependent encryption and decryption functions as $E_K, D_K : \mathbb{F}_2^n \to \mathbb{F}_2^n$. The simplest way to characterize this bijection is through a lookup table which yields the highest possible performance as each block can be encrypted by one simple lookup depending on the key and the message. This is not practical though due to most ciphers working with block and key sizes $n, |K| \geq 64$. For a block cipher with $n = 64, |K| = 128$, a space of $2^{64}2^{128}64 = 2^{198}$ is necessary. Considering modern consumer hard

disks being able to store data in the order of 2^{40}, it is easy to see that a lookup table is wholly impractical. We therefore describe block ciphers algorithmically which opens up possibilities for different tradeoffs and security concerns.

1.2.1 GIFT

GIFT[1], first presented in the *CHES 2017* cryptographic hardware and embedded systems conference, is a lightweight block cipher based on a previous design called PRESENT, developed in 2007. Its goal is to offer maximum security while being extremely light on resources. Modern battery-powered devices like RFID tags or low-latency operations like on-the-fly disc encryption present strong hardware and power constraints. GIFT aims to be a simple, low-energy cipher suited for these kinds of applications.

GIFT comes in two variants; GIFT-64 working with 64-bit blocks and GIFT-128 working with 128-bit blocks. In both cases, the key is 128 bits long. The design is a very simple, round-based substitution-permutation network (SPN). One round consists in a sequential application of the confusion layer by means of 4-bit S-boxes and subsequent diffusion through bit permutation. After the bit permutation, a round key is added to the cipher state and the single round is complete. GIFT-64 uses 28 rounds while GIFT-128 uses 40 rounds.
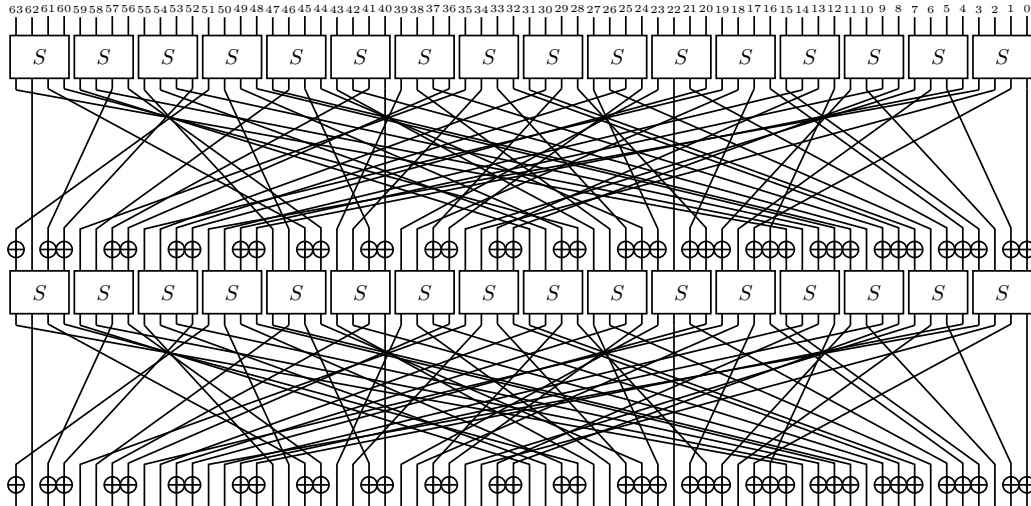

Figure 1.1: Two rounds of GIFT-64

5

Substitution layer

The input of GIFT is split into 4-bit nibbles which are then fed into 16 S-boxes for GIFT-64 and 32 S-boxes for GIFT-128. The S-box $S : \mathbb{F}_2^4 \rightarrow \mathbb{F}_2^4$ is defined as follows:

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
$S(x)$	1	a	4	c	6	f	3	9	2	d	b	7	5	0	8	e

Permutation layer

The permutation P works on individual bits and maps bit b_i to $b_{P(i)}, i \in \{0, 1, \ldots, n-1\}$. The different permutations for GIFT-64 and GIFT-128 can be expressed by:

$$P_{64}(i) = 4 \left\lfloor \frac{i}{16} \right\rfloor + 16 \left(\left(3 \left\lfloor \frac{i \bmod 16}{4} \right\rfloor + (i \bmod 4) \right) \bmod 4 \right) + (i \bmod 4)$$

$$P_{128}(i) = 4 \left\lfloor \frac{i}{16} \right\rfloor + 32 \left(\left(3 \left\lfloor \frac{i \bmod 16}{4} \right\rfloor + (i \bmod 4) \right) \bmod 4 \right) + (i \bmod 4)$$

Round key addition

The last step of each round consists in XORing a round key R_i to the cipher state. The new cipher state s_{i+1} after each full round is therefore given by

$$s_{i+1} = P(S(s_i)) \oplus R_i$$

Round key extraction and key schedule

Round key extraction differs for GIFT-64 and GIFT-128. Let $K = k7||k6|| \ldots ||k0$ denote the 128-bit key state.

GIFT-64 . We extract two 16-bit words $U||V = k_1||k_0$ from the key state. u_i and v_i are XORed to r_{4i+1} and r_{4i} of the round key R respectively.

GIFT-128 . We extract two 32-bit words $U||V = k_5||k4||k1||k_0$ from the key state. u_i and v_i are XORed to r_{4i+2} and b_{4i+1} of the round key R respectively.

In both cases, we additionally XOR a round constant $C = c_5c_4c_3c_2c_1c_0$ to bit positions $n - 1, 23, 19, 15, 11, 7, 3$. The round constants are generated using a 6-bit affine linear-feedback shift register and have the following values:

Rounds	Constants
1 - 16	01,03,07,0F,1F,3E,3D,3B,37,2F,1E,3C,39,33,27,0E
17 - 32	1D,3A,35,2B,16,2C,18,30,21,02,05,0B,17,2E,1C,38
33 - 48	31,23,06,0D,1B,36,2D,1A,34,29,12,24,08,11,22,04

The key state is then updated by setting $k_1 \leftarrow k_1 \ggg 2$, $k_0 \leftarrow k_0 \ggg 12$ and rotating the new state 32 bits to the right:

$$k_7||k_6||\dots||k_1||k_0 \leftarrow k_1 \ggg 2||k_0 \ggg 12||k_7||k_6||\dots||k_3||k_2$$

1.2.2 Camellia

1.3 The ARMv8 platform

With small devices, embedded processors and ASICs becoming ever more ubiquitous and essential in areas like medicine or automotive design, the need for ...

Chapter 2

Implementation strategies

Due to the structural differences of SPN- and Feistel network-based ciphers, we shall analyze these two separately.

2.1 Strategies for SPN

Three implementation strategies for substitution-permutation networks are introduced by [2]:

- Table-based implementations

- `vperm` implementations

- Bitslice implementations

2.1.1 Table-based

Table-driven programming is a simple way to increase performance of operations by tabulating the results, therefore requiring only a single memory access to acquire the result. This approach is obviously limited to manageable table sizes, so while tabulating a function like the AES S-box $S_{AES} : \mathbb{F}_2^8 \to \mathbb{F}_2^8$ requires only 2^{11} space, tabulating the `GIFT` permutation layer $P_{GIFT} : \mathbb{F}_2^{64} \to \mathbb{F}_2^{64}$ would require 2^{70} space, which is totally unfeasible.

A common approach is to tabulate the output of each S-box, including the diffusion layer, and then XORing the results together. Let n denote the internal cipher state size and s the size of a single S-box in bits. For

each S-box $S_i, i \in \{0, \ldots, \frac{n}{s}\}$, we can construct a mapping $T_i : \mathbb{F}_2^s \to \mathbb{F}_2^n$ representing substitution with subsequent permutation of that single S-box. The cipher state before round key addition is then given by $\bigoplus_{i=0}^{\frac{n}{s}-1} T_i(m_i)$ for each s-bit message chunk m_i. This approach requires space of $\frac{n}{s}|\mathbb{F}_2^s|n = \frac{n^2 2^s}{s}$ bits, which, for GIFT-64, results in a manageable size of $\frac{64^2 2^4}{4} = 2^{14}$ bits which equals 16 KiB.

Constructing the tables

For GIFT-64, table construction is relatively straightforward and can be done as follows:

Listing 2.1: Table construction algorithm

```
1    tables <- [][]
2    for sbox_index from 0 to 15 do
3        for sbox_input from 0 to 15 do
4            output <- sbox(sbox_input)
5            output <- permute(output << (4 * sbox_index))
6            tables[sbox_index][sbox_input] <- output
```

Implementing this algorithm gives us the following table representing the first and second S-box.

x	$T_0(x)$	$T_1(x)$	$\ldots$
$0x0$	$0x1$	$0x1000000000000$	$\ldots$
$0x1$	$0x8000000020000$	$0x800000002$	$\ldots$
$0x2$	$0x400000000$	$0x40000$	$\ldots$
$0x3$	$0x8000400000000$	$0x800040000$	$\ldots$
$0x4$	$0x400020000$	$0x40002$	$\ldots$
$0x5$	$0x8000400020001$	$0x1000800040002$	$\ldots$
$0x6$	$0x20001$	$0x1000000000002$	$\ldots$
$0x7$	$0x8000000000001$	$0x1000800000000$	$\ldots$
$0x8$	$0x20000$	$0x2$	$\ldots$
$0x9$	$0x8000400000001$	$0x1000800040000$	$\ldots$
$0xa$	$0x8000000020001$	$0x1000800000002$	$\ldots$
$0xb$	$0x400020001$	$0x1000000040002$	$\ldots$
$0xc$	$0x400000001$	$0x1000000040000$	$\ldots$
$0xd$	$0x0$	$0x0$	$\ldots$
$0xe$	$0x8000000000000$	$0x800000000$	$\ldots$
$0xf$	$0x8000400020000$	$0x800040002$	$\ldots$

The tables for `GIFT-128` can be generated in a similar way by looping through all 32 S-boxes instead of 16 on line 3.

2.1.2   Using `vperm`

Nowadays, most instructions set architectures support single-instruction, multiple-data processing. The idea of such an SIMD system is to work on multiple data stored in vectors at once to speed up calculations. For A64, two types of vector processing are available:

1. Advanced SIMD, known as NEON

2. Scalable Vector Extension (SVE)

We will take a look at NEON as this is the type of vector processing supported by the Cortex-A73 processor.

ARM Neon

The register file of the NEON unit is made up of 32 quad-word (128-bit) registers `V[0-31]`, each extending the standard 64-bit floating-point registers `D[0-31]`. These registers are divided into equally sized lanes on which the vector instructions operate. Valid ways to interpret for example the register `V0` are:
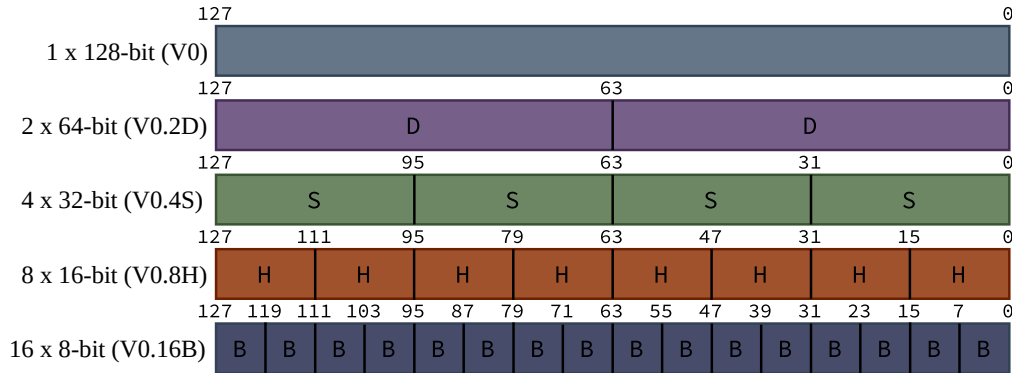

Figure 2.1: Divisions of the V0 register

NEON instructions interpret their operands' layouts (i.e. lane count and width) through the use of suffixes such as `.4S` or `.8H`. For example, adding

eight 16-bit halfwords from register V1 and V2 together and storing the result in V0 can be done as follows:

$$\text{ADD V0.8H, V1.8H, V2.8H}$$

Figure 2.2: Addition of two vector registers

The plenitude of different processing instructions allow flexible ways to further speed up algorithms having reached their optimizational limit on non-SIMD platforms. vperm, a general term standing for *vector permute*, is a common instruction on SIMD machines. Called TBL on NEON, it is used for parallel table lookups and arbitrary permutations. It takes two inputs to perform a lanewise lookup:

1. A register with lookup values

2. Two or more registers containing data

S-box lookup

This instruction can be used to implement S-box lookup of all 16 S-boxes in a single instruction. We do this by packing our 64-bit cipher state $s = s_{15}||s_{14}||\dots||s_0$ into a vector register V_0. Because we can only operate on whole bytes, we put each 4-bit S-box into an 8-bit lane which neatly fits into the 128-bit registers. We then put the S-box itself into register V_1 which will be used as the data register for the table lookup.

The confusion layer can now be performed through one TBL instruction:

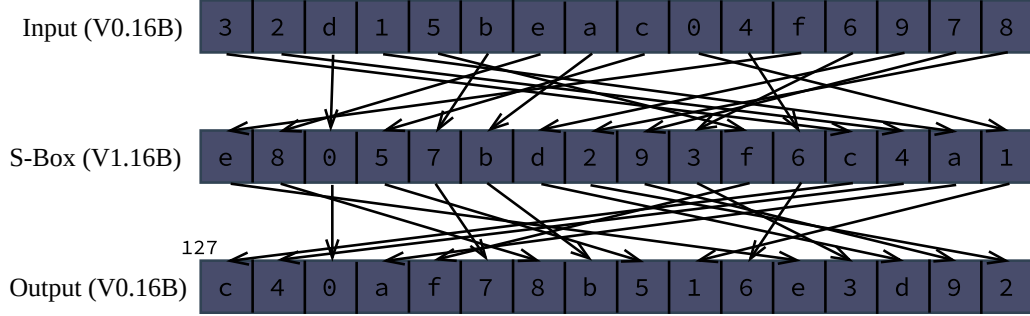$$\text{TBL V0.16B, V1.16B, V0.16B}$$

Figure 2.3: Performing the S-Box lookup in parallel

2.1.3 Bitslicing

Bitslicing refers to the technique of splitting up n bits into m slices to achieve a more efficient representation to operate on. The structure of `GIFT` naturally offers possibilities for bitslicing. We split the cipher state bits $b_{63}b_{62}\ldots b_0$ into four slices $S_i, i \in \{0, 1, 2, 3\}$ such that the i-th slice contains all i-th bits of the individual S-boxes. This is equivalent to transposing the bit matrix.

$$S = \begin{bmatrix} S_0 \\ S_1 \\ S_2 \\ S_3 \end{bmatrix} = \begin{bmatrix} b_{60}b_{56}b_{52}\ldots b_0 \\ b_{61}b_{57}b_{53}\ldots b_1 \\ b_{62}b_{58}b_{54}\ldots b_2 \\ b_{63}b_{59}b_{55}\ldots b_3 \end{bmatrix}$$

Parallel S-Boxes

This representation offers multiple advantages. We first note that computation of the S-box can be executed in parallel, similar to the `vperm` technique above. This can be done by finding an algorithmic way to apply the S-box which has already been proposed by the original `GIFT` authors:

$$S_1 \leftarrow S_1 \oplus (S_0 \wedge S_2)$$
$$t \leftarrow S_0 \oplus (S_1 \wedge S_3)$$
$$S_2 \leftarrow S_2 \oplus (t \vee S_1)$$
$$S_0 \leftarrow S_3 \oplus S_2$$
$$S_1 \leftarrow S_1 \oplus S_0$$
$$S_0 \leftarrow \neg S_0$$
$$S_2 \leftarrow S_2 \oplus (t \wedge S_1)$$
$$S_3 \leftarrow t$$

This is very efficient as it only requires six XOR-, three AND and one OR operation.

An important property of the permutation is the fact that bits always stay in their slice. This means we can decompose the permutation P into four permutations $P_i, i \in \{0, 1, 2, 3\}$ and apply these permutations separately to each slice. One possible way to implement a permutation P_i in software is to mask off all bits individually, shift them to their correct position and OR them together:

$$P_i(S_i) = \bigvee_{k=0}^{15} (S_i \wedge m_i) \lll s_i$$

This approach requires 47 operations, meaning all four permutations require over 150 operations which would present a major bottleneck to the round function. We can improve on this by working on multiple message blocks at once and using the aforementioned `vperm` instruction to implement the bit shuffling. We then need only four instructions for the complete diffusion layer.

Using `vperm` for slice permutation

We cannot use the `TBL` instruction directly as we need to shuffle individual bits, but the smallest data we can operate on are bytes. We therefore encrypt $8n$ messages at once which allows us to create bytewise groupings. These messages are put into $4m$ registers with register R_{4i} containing S_0, register R_{4i+1} containing S_1 and so forth. With block size BS and register size RS, the following must hold:

$$8n \cdot BS = 4m \cdot RS$$

In the case of `GIFT-64` with $BS = 64$ and ARM NEON with $RS = 128$, we get

$$8n \cdot 64 = 4m \cdot 128 \Leftrightarrow n = m$$

$n = m = 1$ would be a valid choice which yields eight messages divided into four registers. We choose $n = m = 2$ so we can directly utilize the algorithm for bit packing presented by the original GIFT authors, although it is simple to adapt this algorithm to only four registers and eight messages by adjusting the `SWAPMOVE` shift and mask values.

Packing the data into bitslice format

Let $a, b, \ldots, p$ be sixteen messages of length 64 with subscripts denoting individual bits. We first put these messages into eight SIMD registers $V_0, V_1, \ldots, V_7$:

$$
\begin{aligned}
V_0 &= b||a & V_4 &= j||i \\
V_1 &= d||c & V_5 &= l||\,k \\
V_2 &= f||e & V_6 &= n||m \\
V_3 &= h||g & V_7 &= p||o
\end{aligned}
$$

We then use the `SWAPMOVE` technique to bring the data into bitslice format. This operation operates on two registers A, B using mask M and shift value N. It swaps bits in A masked by $(M \ll N)$ with bits in B masked by M in using only three XOR-, one AND- and two shift operations.

$$
\begin{aligned}
&\text{SWAPMOVE}(A, B, M, N): \\
&\quad T = ((A \gg N) \oplus B) \wedge M \\
&\quad B = B \oplus T \\
&\quad A = A \oplus (T \ll N)
\end{aligned}
$$

One caveat of this approach is the fact that NEON registers cannot be shifted in their entirety due to the fact bits are not able to cross lanes. This

leads to the problem of being able to shift at most two lanes of 64 bits at once. We thus need to implement the `shr(V,n)` and `shl(V,n)` operations on our own. This can be done by first extracting the 64-bit lanes a, b out of $V = b||a$, shifting the lanes individually and finally shifting and ORing the crossing bits back into the other lane.

$$
\begin{aligned}
\text{shl}(V, n) : & \\
a, b \quad &= V[0], V[1] \\
c \quad &= (a \gg (64 - n)) \\
a \quad &= (a \ll n) \\
b \quad &= (b \ll n) \vee c \\
V[0], V[1] &= a, b
\end{aligned}
$$

The following operations group all ith bits of the messages $a, c, \ldots, o$ into bytes and puts these into the lower half of the registers $V_{i \bmod 8}$. The same is done for messages $b, d, \ldots, p$, only differing in that the bytes are put into the upper half of the registers.

$\text{SWAPMOVE}(V_0, V_1, 0x5555\ldots55, 1)$ $\quad$ $\text{SWAPMOVE}(V_4, V_5, 0x5555\ldots55, 1)$

$\text{SWAPMOVE}(V_2, V_3, 0x5555\ldots55, 1)$ $\quad$ $\text{SWAPMOVE}(V_6, V_7, 0x5555\ldots55, 1)$

$\text{SWAPMOVE}(V_0, V_2, 0x3333\ldots33, 2)$ $\quad$ $\text{SWAPMOVE}(V_4, V_6, 0x3333\ldots33, 2)$

$\text{SWAPMOVE}(V_1, V_3, 0x3333\ldots33, 2)$ $\quad$ $\text{SWAPMOVE}(V_5, V_7, 0x3333\ldots33, 2)$

$\text{SWAPMOVE}(V_0, V_4, 0x0f0f\ldots0f, 4)$ $\quad$ $\text{SWAPMOVE}(V_1, V_5, 0x0f0f\ldots0f, 4)$

$\text{SWAPMOVE}(V_2, V_6, 0x0f0f\ldots0f, 4)$ $\quad$ $\text{SWAPMOVE}(V_3, V_7, 0x0f0f\ldots0f, 4)$

With $Ax = o_x m_x k_x j_x g_x e_x c_x a_x$ and $Bx = p_x n_x l_x i_x h_x f_x d_x b_x$ denoting byte groups, our data now has the following permutation-friendly format:

n	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
V_0	$B56$	$B48$	$B40$	$B32$	$B24$	$B16$	$B8$	$B0$	$A56$	$A48$	$A40$	$A32$	$A24$	$A16$	$A8$	$A0$
V_1	$B57$	$B49$	$B41$	$B33$	$B25$	$B17$	$B9$	$B1$	$A57$	$A49$	$A41$	$A33$	$A25$	$A17$	$A9$	$A1$
V_2	$B58$	$B50$	$B42$	$B34$	$B26$	$B18$	$B10$	$B2$	$A58$	$A50$	$A42$	$A34$	$A26$	$A18$	$A10$	$A2$
V_3	$B59$	$B51$	$B43$	$B35$	$B27$	$B19$	$B11$	$B3$	$A59$	$A51$	$A43$	$A35$	$A27$	$A19$	$A11$	$A3$
V_4	$B60$	$B52$	$B44$	$B36$	$B28$	$B20$	$B12$	$B4$	$A60$	$A52$	$A44$	$A36$	$A28$	$A20$	$A12$	$A4$
V_5	$B61$	$B53$	$B45$	$B37$	$B29$	$B21$	$B13$	$B5$	$A61$	$A53$	$A45$	$A37$	$A29$	$A21$	$A13$	$A5$
V_6	$B62$	$B54$	$B46$	$B38$	$B30$	$B22$	$B14$	$B6$	$A62$	$A54$	$A46$	$A38$	$A30$	$A22$	$A14$	$A6$
V_7	$B63$	$B55$	$B47$	$B39$	$B31$	$B23$	$B15$	$B7$	$A63$	$A55$	$A47$	$A39$	$A31$	$A23$	$A15$	$A7$

Although this would already work, we prefer to have only bits of the same messages in each register - otherwise the permutation would need to operate on two source registers with the added requirement of storing the pre-permutation values for the first four registers, slowing down the round function through superfluous load/stores. This transformation is trivial by use of TBL with two data source operands. The final data format we operate on is as follows:

n	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
V_0	$A60$	$A56$	$A52$	$A48$	$A44$	$A40$	$A36$	$A32$	$A28$	$A24$	$A20$	$A16$	$A12$	$A8$	$A4$	$A0$
V_1	$A61$	$A57$	$A53$	$A49$	$A45$	$A41$	$A37$	$A33$	$A29$	$A25$	$A21$	$A17$	$A13$	$A9$	$A5$	$A1$
V_2	$A62$	$A58$	$A54$	$A50$	$A46$	$A42$	$A38$	$A34$	$A30$	$A26$	$A22$	$A18$	$A14$	$A10$	$A6$	$A2$
V_3	$A63$	$A59$	$A55$	$A51$	$A47$	$A43$	$A39$	$A35$	$A31$	$A27$	$A23$	$A19$	$A15$	$A11$	$A7$	$A3$
V_4	$B60$	$B56$	$B52$	$B48$	$B44$	$B40$	$B36$	$B32$	$B28$	$B24$	$B20$	$B16$	$B12$	$B8$	$B4$	$B0$
V_5	$B61$	$B57$	$B53$	$B49$	$B45$	$B41$	$B37$	$B33$	$B29$	$B25$	$B21$	$B17$	$B13$	$B9$	$B5$	$B1$
V_6	$B62$	$B58$	$B54$	$B50$	$B46$	$B42$	$B38$	$B34$	$B30$	$B26$	$B22$	$B18$	$B14$	$B10$	$B6$	$B2$
V_7	$B63$	$B59$	$B55$	$B51$	$B47$	$B43$	$B39$	$B35$	$B31$	$B27$	$B23$	$B19$	$B15$	$B11$	$B7$	$B3$

We can now create permutation tables using the specification of the individual slice permutations P_i which are then applied to V_i and V_{i+4}:

j	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
$P_0(j)$	0	12	8	4	1	13	9	5	2	14	10	6	3	15	11	7
$P_1(j)$	4	0	12	8	5	1	13	9	6	2	14	10	7	3	15	11
$P_2(j)$	8	4	0	12	9	5	1	13	10	6	2	14	11	7	3	15
$P_3(j)$	12	8	4	0	13	9	5	1	14	10	6	2	15	11	7	3

One thing to take note of is the original permutation values only show where a given byte should land, not which byte belongs to a certain position - i.e. for P_0, byte 1 should land in position 12, but the byte belonging to position 1 is byte 4. Because TBL works in the latter way, we have to do some trivial rearrangements.

Assuming the correct permutation values are put into registers V_8, V_9, V_{10}, V_{11}, this now allows us to compute the permutation layer for all 16 blocks in only eight permutation instructions.

```
TBL V0, V0, V8      TBL V1, V1, V9
TBL V4, V4, V8      TBL V5, V5, V9
TBL V2, V2, V10     TBL V3, V3, V11
TBL V6, V6, V10     TBL V7, V7, V11
```

Round key function

In contrast to packing and unpacking of data which is only done once in the beginning and end, a round key is derived for every round, so the round key derivation function needs to be as fast as possible. A simple but naive approach for one round would be to generate a single round key, copy it 15 times and pack the resulting registers similar to how we proceed with the messages. Due to the cost of packing the messages, this is prohibitivly expensive. Because we know where each byte group ends up after packing, we can directly XOR the round key bits to the correct position. Extending these bits to bytes can then be done simply by repeatedly shifting and ORing the registers together.

2.2 Strategies for Camellia

Chapter 3

Implementation

We will provide and discuss implementations for the presented strategies in the C programming language. Although directly writing Assembler code could result in a small performance benefit, this generally increases the work necessary by an order of magnitude for only limited results. Instruction-level optimization and in particular register allocation is left to the compiler.

3.1   Table-based

We generate a 16 by 16 array of **uint64_t** as explained in 3. The round function then consists of extracting each nibble, looking up the corresponding table value and ORing this to the new cipher state.

Listing 3.1: Table-based round function

```
 1  static const uint64_t table[16][16] = {
 2          { 0x0000000000000001UL, 0x0008000000020000UL,
 3            0x0000000400000000UL, 0x0008000400000000UL,
 4            0x0000000400020000UL, 0x0008000400020001UL,
 5            0x0000000000020001UL, 0x0008000000000001UL,
 6            0x0000000000020000UL, 0x0008000400000001UL,
 7            0x0008000000020001UL, 0x0000000400020001UL,
 8            0x0000000400000001UL, 0x0000000000000000UL,
 9            0x0008000000000000UL, 0x0008000400020000UL },
10          // ...
11          { 0x0000000010000000UL, 0x0000020000008000UL,
12            0x4000000000000000UL, 0x4000000000008000UL,
13            0x4000200000000000UL, 0x4000200010008000UL,
14            0x0000200010000000UL, 0x0000000010008000UL,
15            0x0000020000000000UL, 0x4000000010008000UL,
16            0x0000200010008000UL, 0x4000200010000000UL,
17            0x4000000010000000UL, 0x0000000000000000UL,
18            0x0000000000008000UL, 0x4000200000008000UL }
```

```
19  };
20
21  uint64_t gift_64_table_subperm(const uint64_t cipher_state)
22  {
23          uint64_t new_cipher_state = 0;
24
25          for (size_t i = 0; i < 16; i++) {
26                  int nibble = (cipher_state >> (i * 4)) & 0xf;
27                  new_cipher_state |= table[i][nibble];
28          }
29
30          return new_cipher_state;
31  }
```

The round function is similar for all strategies: for each round, we extract 16 bits from the key state which are then XORed to the round key, together with the precomputed round constants and the single bit. Key state updating is done using masks and shifts only which gets optimized by the compiler through bit-field extraction instructions which can extract, shift and insert a bit field at once. Intermediate values are also optimized away.

Listing 3.2: Round key generation function

```
1   void gift_64_table_generate_round_keys(uint64_t round_keys[restrict
        ROUNDS_GIFT_64],
2                                          const uint64_t key[restrict 2])
3   {
4           uint64_t key_state[] = {key[0], key[1]};
5           for (int round = 0; round < ROUNDS_GIFT_64; round++) {
6                   int v = (key_state[0] >> 0 ) & 0xffff;
7                   int u = (key_state[0] >> 16) & 0xffff;
8
9                   // add round key (RK=U||V)
10                  round_keys[round] = 0UL;
11                  for (size_t i = 0; i < 16; i++) {
12                          int key_bit_v  = (v >> i)  & 0x1;
13                          int key_bit_u  = (u >> i)  & 0x1;
14                          round_keys[round] ^= (uint64_t)key_bit_v << (i * 4 + 0);
15                          round_keys[round] ^= (uint64_t)key_bit_u << (i * 4 + 1);
16                  }
17
18                  // add single bit
19                  round_keys[round] ^= 1UL << 63;
20
21                  // add round constants
22                  round_keys[round] ^= ((round_constant[round] >> 0) & 0x1) << 3;
23                  round_keys[round] ^= ((round_constant[round] >> 1) & 0x1) << 7;
24                  round_keys[round] ^= ((round_constant[round] >> 2) & 0x1) << 11;
25                  round_keys[round] ^= ((round_constant[round] >> 3) & 0x1) << 15;
26                  round_keys[round] ^= ((round_constant[round] >> 4) & 0x1) << 19;
27                  round_keys[round] ^= ((round_constant[round] >> 5) & 0x1) << 23;
28
29                  // update key state
30                  int k0 = (key_state[0] >> 0 ) & 0xffffUL;
```

```
31              int k1 = (key_state[0] >> 16) & 0xffffUL;
32              k0 = (k0 >> 12) | ((k0 & 0xfff) << 4);
33              k1 = (k1 >> 2 ) | ((k1 & 0x3  ) << 14);
34              key_state[0] >>= 32;
35              key_state[0] |= (key_state[1] & 0xffffffffUL) << 32;
36              key_state[1] >>= 32;
37              key_state[1] |= ((uint64_t)k0 << 32) | ((uint64_t)k1 << 48);
38          }
39  }
```

While this approach does not rely on ARM-specific instructions, the next two strategies do. This is why we use NEON intrinsics.

3.2   Using **vperm**

3.2.1   NEON intrinsics

The header file <arm_neon.h> provides ARM-specific data and function definitions including vector data types and C functions for working with these vectors. These functions are known as NEON intrinsics and give the programmer a high-level interface to most NEON instructions. Major advantages of this approach include the ease of development as the compiler takes over register allocation and load/store operations as well as performance benefits through compiler optimizations.

Standard vector data types have the format uintnxm_t with lane width n in bits and and lane count m. Array types of the format uintnxmxc_t, $c \in \{2, 3, 4\}$ are also defined which are used in operations requiring multiple parameters like TBL or pairwise load/stores. Intrinsics start with v, optionally contain a q after the operation name to indicate operation on a 128-bit register, and end with the lane data format. Multiplying eight pairs of 16-bit numbers a,b for example can be done via the following:

uint16x8_t result = vmulq_u16(a, b);

In this case, the compiler allocates vector registers for a, b and result and assembles the intrinsic to MUL Vr.8H, Va.8H, Vb.8H. Necessary loads and stores for the result and parameters are also handled automatically. Of special interest to us are the following intrinsics, each existing in different variants with different lane widths and also array types:

Intrinsic		Description
`uint8x16_t`	`vreinterpretq_u8_u64(uint64x2_t)`	Explicit casting
`uint64_t`	`vgetq_lane_u64(void)`	Extract a single lane
`void`	`vsetq_lane_u64(uint64_t)`	Insert a single lane
`uint64x2_t`	`vdupq_n_u64(uint64_t)`	Initialize all lanes to same value
`void`	`vst1q_u64(uint64_t*, uint64x2_t)`	Store from register to memory
`uint64x2_t`	`vld1q_u64(uint64_t*, uint64x2_t)`	Load from memory to register
`uint8x16_t`	`veorq_u8(uint8x16_t, uint8x16_t)`	bitwise XOR
`uint8x16_t`	`vandq_u8(uint8x16_t, uint8x16_t)`	bitwise AND
`uint8x16_t`	`vorrq_u8(uint8x16_t, uint8x16_t)`	bitwise OR
`uint8x16_t`	`vmvnq_u8(uint8x16_t)`	bitwise NOT
`uint8x16_t`	`vqtbl2q_u8(uint8x16_t, uint8x16_t)`	permutation (`TBL`)

3.3   Bitslicing

Chapter 4

Evaluation

Acknowledgements

I want to thank ...

Chapter 5

Appendix

Listing 5.1: gift_vec_sliced.h

```
 1  #ifndef GIFT_VEC_SLICED_H
 2  #define GIFT_VEC_SLICED_H
 3
 4  #include <stdint.h>
 5  #include <arm_neon.h>
 6
 7  #define ROUNDS_GIFT_64 28
 8
 9  // expose for benchmarking
10  uint8x16_t shl(uint8x16_t v, int n);
11  uint8x16_t shr(uint8x16_t v, int n);
12  void gift_64_vec_sliced_swapmove(uint8x16_t *restrict a, uint8x16_t *restrict b,
13                                   uint8x16_t m, int n);
14  void gift_64_vec_sliced_bits_pack(uint8x16x4_t m[restrict 2]);
15
16  void gift_64_vec_sliced_subcells(uint8x16x4_t cipher_state[restrict 2]);
17  void gift_64_vec_sliced_subcells_inv(uint8x16x4_t cipher_state[restrict 2]);
18  void gift_64_vec_sliced_permute(uint8x16x4_t cipher_state[restrict 2]);
19  void gift_64_vec_sliced_permute_inv(uint8x16x4_t cipher_state[2]);
20  void gift_64_vec_sliced_generate_round_keys(uint8x16x4_t round_keys[restrict
        ROUNDS_GIFT_64][2],
21                                              const uint64_t key[restrict 2]);
22
23  void gift_64_vec_sliced_init(void);
24
25  void gift_64_vec_sliced_encrypt(uint64_t c[restrict 16],
26                                  const uint64_t m[restrict 16],
27                                  const uint64_t key[restrict 2]);
28  void gift_64_vec_sliced_decrypt(uint64_t m[restrict 16],
29                                  const uint64_t c[restrict 16],
30                                  const uint64_t key[restrict 2]);
31
32  #endif
```

Listing 5.2: gift_vec_sliced.c

```c
1   #include "gift_vec_sliced.h"
2
3   #include <stddef.h>
4   #include <stdio.h>
5   #include <string.h>
6
7   static uint64_t perm_u64[16] = {
8           0x1410050115110400UL, 0x1410050115110400UL,
9           0x0501151104001410UL, 0x0501151104001410UL,
10          0x1511040014100501UL, 0x1511040014100501UL,
11          0x0400141005011511UL, 0x0400141005011511UL,
12          0x1612070317130602UL, 0x1612070317130602UL,
13          0x0703171306021612UL, 0x0703171306021612UL,
14          0x1713060216120703UL, 0x1713060216120703UL,
15          0x0602161207031713UL, 0x0602161207031713UL
16  };
17
18  static uint64_t perm_inv_u64[16] = {
19          0x1511050114100400UL, 0x1511050114100400UL,
20          0x1713070316120602UL, 0x1713070316120602UL,
21          0x1115010510140004UL, 0x1115010510140004UL,
22          0x1317030712160206UL, 0x1317030712160206UL,
23          0x1317030712160206UL, 0x1317030712160206UL,
24          0x1511050114100400UL, 0x1511050114100400UL,
25          0x1713070316120602UL, 0x1713070316120602UL,
26          0x1115010510140004UL, 0x1115010510140004UL
27  };
28
29  static uint8x16x4_t perm[2];
30  static uint8x16x4_t perm_inv[2];
31
32  static uint8x16_t pack_mask_0;
33  static uint8x16_t pack_mask_1;
34  static uint8x16_t pack_mask_2;
35
36  static const int round_constant[] = {
37          // rounds 0-15
38          0x01, 0x03, 0x07, 0x0F, 0x1F, 0x3E, 0x3D, 0x3B,
39          0x37, 0x2F, 0x1E, 0x3C, 0x39, 0x33, 0x27, 0x0E,
40          // rounds 16-31
41          0x1D, 0x3A, 0x35, 0x2B, 0x16, 0x2C, 0x18, 0x30,
42          0x21, 0x02, 0x05, 0x0B, 0x17, 0x2E, 0x1C, 0x38,
43          // rounds 32-47
44          0x31, 0x23, 0x06, 0x0D, 0x1B, 0x36, 0x2D, 0x1A,
45          0x34, 0x29, 0x12, 0x24, 0x08, 0x11, 0x22, 0x04
46  };
47
48  uint8x16_t shl(uint8x16_t v, int n)
49  {
50          uint64_t l[2];
51          vst1q_u64(l, v);
52          l[1] = l[1] << n | (l[0] >> (64 - n));
53          l[1] <<= n;
54          return vreinterpretq_u8_u64(vld1q_u64(l));
55  }
56
```

```
57  uint8x16_t shr(uint8x16_t v, int n)
58  {
59          uint64_t l[2];
60          vst1q_u64(l, v);
61          l[0] = l[0] >> n | (((l[1] << (64 - n)) >> (64 - n)) << (64 - n));
62          l[1] >>= n;
63          return vreinterpretq_u8_u64(vld1q_u64(l));
64  }
65
66  void gift_64_vec_sliced_swapmove(uint8x16_t *restrict a, uint8x16_t *restrict b,
67          uint8x16_t m, int n)
67  {
68
69          uint8x16_t t = vandq_u8(veorq_u8(shr(*a, n), *b), m);
70          *b = veorq_u8(*b, t);
71          *a = veorq_u8(*a, shl(t, n));
72  }
73
74  void gift_64_vec_sliced_bits_pack(uint8x16x4_t m[restrict 2])
75  {
76          // take care not to shift mask bits out of the register
77          gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[1], pack_mask_0, 1);
78          gift_64_vec_sliced_swapmove(&m[0].val[2], &m[0].val[3], pack_mask_0, 1);
79          gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[1], pack_mask_0, 1);
80          gift_64_vec_sliced_swapmove(&m[1].val[2], &m[1].val[3], pack_mask_0, 1);
81
82          gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[2], pack_mask_1, 2);
83          gift_64_vec_sliced_swapmove(&m[0].val[1], &m[0].val[3], pack_mask_1, 2);
84          gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[2], pack_mask_1, 2);
85          gift_64_vec_sliced_swapmove(&m[1].val[1], &m[1].val[3], pack_mask_1, 2);
86
87          // make bytes (a0 b0 c0 d0 a4 b4 c4 d4 -> a0 b0 c0 d0 e0 f0 g0 h0)
88          gift_64_vec_sliced_swapmove(&m[0].val[0], &m[1].val[0], pack_mask_2, 4);
89          gift_64_vec_sliced_swapmove(&m[0].val[2], &m[1].val[2], pack_mask_2, 4);
90          gift_64_vec_sliced_swapmove(&m[0].val[1], &m[1].val[1], pack_mask_2, 4);
91          gift_64_vec_sliced_swapmove(&m[0].val[3], &m[1].val[3], pack_mask_2, 4);
92  }
93
94  void gift_64_vec_sliced_subcells(uint8x16x4_t cs[restrict 2])
95  {
96          cs[0].val[1] = veorq_u8(cs[0].val[1], vandq_u8(cs[0].val[0], cs[0].val
                  [2]));
97          uint8x16_t t = veorq_u8(cs[0].val[0], vandq_u8(cs[0].val[1], cs[0].val
                  [3]));
98          cs[0].val[2] = veorq_u8(cs[0].val[2], vorrq_u8(t, cs[0].val[1]));
99          cs[0].val[0] = veorq_u8(cs[0].val[3], cs[0].val[2]);
100         cs[0].val[1] = veorq_u8(cs[0].val[1], cs[0].val[0]);
101         cs[0].val[0] = vmvnq_u8(cs[0].val[0]);
102         cs[0].val[2] = veorq_u8(cs[0].val[2], vandq_u8(t, cs[0].val[1]));
103         cs[0].val[3] = t;
104
105         cs[1].val[1] = veorq_u8(cs[1].val[1], vandq_u8(cs[1].val[0], cs[1].val
                  [2]));
106         t            = veorq_u8(cs[1].val[0], vandq_u8(cs[1].val[1], cs[1].val
                  [3]));
107         cs[1].val[2] = veorq_u8(cs[1].val[2], vorrq_u8(t, cs[1].val[1]));
108         cs[1].val[0] = veorq_u8(cs[1].val[3], cs[1].val[2]);
```

```
109            cs[1].val[1] = veorq_u8(cs[1].val[1], cs[1].val[0]);
110            cs[1].val[0] = vmvnq_u8(cs[1].val[0]);
111            cs[1].val[2] = veorq_u8(cs[1].val[2], vandq_u8(t, cs[1].val[1]));
112            cs[1].val[3] = t;
113    }
114
115    void gift_64_vec_sliced_subcells_inv(uint8x16x4_t cs[restrict 2])
116    {
117            uint8x16_t t = cs[0].val[3];
118            cs[0].val[2] = veorq_u8(cs[0].val[2], vandq_u8(t, cs[0].val[1]));
119            cs[0].val[0] = vmvnq_u8(cs[0].val[0]);
120            cs[0].val[1] = veorq_u8(cs[0].val[1], cs[0].val[0]);
121            cs[0].val[3] = veorq_u8(cs[0].val[0], cs[0].val[2]);
122            cs[0].val[2] = veorq_u8(cs[0].val[2], vorrq_u8(t, cs[0].val[1]));
123            cs[0].val[0] = veorq_u8(t, vandq_u8(cs[0].val[1], cs[0].val[3]));
124            cs[0].val[1] = veorq_u8(cs[0].val[1], vandq_u8(cs[0].val[0], cs[0].val
                   [2]));
125
126            t            = cs[1].val[3];
127            cs[1].val[2] = veorq_u8(cs[1].val[2], vandq_u8(t, cs[1].val[1]));
128            cs[1].val[0] = vmvnq_u8(cs[1].val[0]);
129            cs[1].val[1] = veorq_u8(cs[1].val[1], cs[1].val[0]);
130            cs[1].val[3] = veorq_u8(cs[1].val[0], cs[1].val[2]);
131            cs[1].val[2] = veorq_u8(cs[1].val[2], vorrq_u8(t, cs[1].val[1]));
132            cs[1].val[0] = veorq_u8(t, vandq_u8(cs[1].val[1], cs[1].val[3]));
133            cs[1].val[1] = veorq_u8(cs[1].val[1], vandq_u8(cs[1].val[0], cs[1].val
                   [2]));
134    }
135
136    void gift_64_vec_sliced_permute(uint8x16x4_t cs[restrict 2])
137    {
138            uint8x16x2_t pairs[4] = {
139                    { .val = { cs[0].val[0], cs[1].val[0] }},
140                    { .val = { cs[0].val[1], cs[1].val[1] }},
141                    { .val = { cs[0].val[2], cs[1].val[2] }},
142                    { .val = { cs[0].val[3], cs[1].val[3] }},
143            };
144
145            cs[0].val[0] = vqtbl2q_u8(pairs[0], perm[0].val[0]);
146            cs[0].val[1] = vqtbl2q_u8(pairs[1], perm[0].val[1]);
147            cs[0].val[2] = vqtbl2q_u8(pairs[2], perm[0].val[2]);
148            cs[0].val[3] = vqtbl2q_u8(pairs[3], perm[0].val[3]);
149
150            cs[1].val[0] = vqtbl2q_u8(pairs[0], perm[1].val[0]);
151            cs[1].val[1] = vqtbl2q_u8(pairs[1], perm[1].val[1]);
152            cs[1].val[2] = vqtbl2q_u8(pairs[2], perm[1].val[2]);
153            cs[1].val[3] = vqtbl2q_u8(pairs[3], perm[1].val[3]);
154    }
155
156    void gift_64_vec_sliced_permute_inv(uint8x16x4_t cs[restrict 2])
157    {
158            uint8x16x2_t pairs[4] = {
159                    { .val = { cs[0].val[0], cs[1].val[0] }},
160                    { .val = { cs[0].val[1], cs[1].val[1] }},
161                    { .val = { cs[0].val[2], cs[1].val[2] }},
162                    { .val = { cs[0].val[3], cs[1].val[3] }},
163            };
```

```
164
165            cs[0].val[0] = vqtbl2q_u8(pairs[0], perm_inv[0].val[0]);
166            cs[0].val[1] = vqtbl2q_u8(pairs[1], perm_inv[0].val[1]);
167            cs[0].val[2] = vqtbl2q_u8(pairs[2], perm_inv[0].val[2]);
168            cs[0].val[3] = vqtbl2q_u8(pairs[3], perm_inv[0].val[3]);
169
170            cs[1].val[0] = vqtbl2q_u8(pairs[0], perm_inv[1].val[0]);
171            cs[1].val[1] = vqtbl2q_u8(pairs[1], perm_inv[1].val[1]);
172            cs[1].val[2] = vqtbl2q_u8(pairs[2], perm_inv[1].val[2]);
173            cs[1].val[3] = vqtbl2q_u8(pairs[3], perm_inv[1].val[3]);
174 }
175
176 void gift_64_vec_sliced_generate_round_keys(uint8x16x4_t round_keys[restrict
       ROUNDS_GIFT_64][2],
177                                             const uint64_t key[restrict 2])
178 {
179        // TODO shift-and-or
180
181        uint64_t key_state[] = {key[0], key[1]};
182        for (int round = 0; round < ROUNDS_GIFT_64; round++) {
183                int v = (key_state[0] >> 0 ) & 0xffff;
184                int u = (key_state[0] >> 16) & 0xffff;
185
186                // add round key (RK=U||V)
187                uint64_t round_key[8] = { 0x0UL };
188                for (size_t i = 0; i < 8; i++) {
189                        int key_bit_v  = (v >> (2 * i + 0)) & 0x1;
190                        int key_bit_u  = (u >> (2 * i + 0)) & 0x1;
191                        round_key[0] ^= (uint64_t)key_bit_v << (i * 8);
192                        round_key[1] ^= (uint64_t)key_bit_u << (i * 8);
193
194                        key_bit_v  = (v >> (2 * i + 1)) & 0x1;
195                        key_bit_u  = (u >> (2 * i + 1)) & 0x1;
196                        round_key[4] ^= (uint64_t)key_bit_v << (i * 8);
197                        round_key[5] ^= (uint64_t)key_bit_u << (i * 8);
198                }
199
200                // add single bit
201                round_key[7] ^= 1UL << (7 * 8);
202
203                // add round constants
204                round_key[3] ^= ((round_constant[round] >> 0) & 0x1) << (0 * 8);
205                round_key[7] ^= ((round_constant[round] >> 1) & 0x1) << (0 * 8);
206                round_key[3] ^= ((round_constant[round] >> 2) & 0x1) << (1 * 8);
207                round_key[7] ^= ((round_constant[round] >> 3) & 0x1) << (1 * 8);
208                round_key[3] ^= ((round_constant[round] >> 4) & 0x1) << (2 * 8);
209                round_key[7] ^= ((round_constant[round] >> 5) & 0x1) << (2 * 8);
210
211                // extend bits to bytes
212                for (size_t i = 0; i < 8; i++) {
213                        round_key[i] |= round_key[i] << 1;
214                        round_key[i] |= round_key[i] << 2;
215                        round_key[i] |= round_key[i] << 4;
216                }
217
218                // load into vector registers (upper and lower part are same)
219                round_keys[round][0].val[0] = vdupq_n_u64(round_key[0]);
```

```
220                        round_keys[round][0].val[1] = vdupq_n_u64(round_key[1]);
221                        round_keys[round][0].val[2] = vdupq_n_u64(round_key[2]);
222                        round_keys[round][0].val[3] = vdupq_n_u64(round_key[3]);
223                        round_keys[round][1].val[0] = vdupq_n_u64(round_key[4]);
224                        round_keys[round][1].val[1] = vdupq_n_u64(round_key[5]);
225                        round_keys[round][1].val[2] = vdupq_n_u64(round_key[6]);
226                        round_keys[round][1].val[3] = vdupq_n_u64(round_key[7]);
227
228                        // update key state
229                        int k0 = (key_state[0] >> 0 ) & 0xffffUL;
230                        int k1 = (key_state[0] >> 16) & 0xffffUL;
231                        k0 = (k0 >> 12) | ((k0 & 0xfff) << 4);
232                        k1 = (k1 >> 2 ) | ((k1 & 0x3  ) << 14);
233                        key_state[0] >>= 32;
234                        key_state[0] |= (key_state[1] & 0xffffffffUL) << 32;
235                        key_state[1] >>= 32;
236                        key_state[1] |= ((uint64_t)k0 << 32) | ((uint64_t)k1 << 48);
237                }
238 }
239
240 void gift_64_vec_sliced_init(void)
241 {
242         // permutations
243         perm[0] = vld1q_u8_x4((uint8_t*)&perm_u64[0]);
244         perm[1] = vld1q_u8_x4((uint8_t*)&perm_u64[8]);
245
246         // inverse permutations
247         perm_inv[0] = vld1q_u8_x4((uint8_t*)&perm_inv_u64[0]);
248         perm_inv[1] = vld1q_u8_x4((uint8_t*)&perm_inv_u64[8]);
249
250         // packing masks
251         pack_mask_0 = vdupq_n_u8(0x55);
252         pack_mask_1 = vdupq_n_u8(0x33);
253         pack_mask_2 = vdupq_n_u8(0x0f);
254 }
255
256 void gift_64_vec_sliced_encrypt(uint64_t c[restrict 16],
257                                 const uint64_t m[restrict 16],
258                                 const uint64_t key[restrict 2])
259 {
260         uint8x16x4_t s[2];
261         s[0] = vld1q_u8_x4((uint8_t*)&m[0]);
262         s[1] = vld1q_u8_x4((uint8_t*)&m[8]);
263         gift_64_vec_sliced_bits_pack(s);
264
265         uint8x16x4_t round_keys[ROUNDS_GIFT_64][2];
266         gift_64_vec_sliced_generate_round_keys(round_keys, key);
267
268         for (int round = 0; round < ROUNDS_GIFT_64; round++) {
269                 gift_64_vec_sliced_subcells(s);
270                 gift_64_vec_sliced_permute(s);
271
272                 s[0].val[0] = veorq_u8(s[0].val[0], round_keys[round][0].val[0]);
273                 s[0].val[1] = veorq_u8(s[0].val[1], round_keys[round][0].val[1]);
274                 s[0].val[2] = veorq_u8(s[0].val[2], round_keys[round][0].val[2]);
275                 s[0].val[3] = veorq_u8(s[0].val[3], round_keys[round][0].val[3]);
276                 s[1].val[0] = veorq_u8(s[1].val[0], round_keys[round][1].val[0]);
```

```
277          s[1].val[1] = veorq_u8(s[1].val[1], round_keys[round][1].val[1]);
278          s[1].val[2] = veorq_u8(s[1].val[2], round_keys[round][1].val[2]);
279          s[1].val[3] = veorq_u8(s[1].val[3], round_keys[round][1].val[3]);
280      }
281
282      gift_64_vec_sliced_bits_pack(s);
283      vst1q_u8_x4((uint8_t*)&c[0], s[0]);
284      vst1q_u8_x4((uint8_t*)&c[8], s[1]);
285  }
286
287  void gift_64_vec_sliced_decrypt(uint64_t m[restrict 16],
288                                  const uint64_t c[restrict 16],
289                                  const uint64_t key[restrict 2])
290  {
291      uint8x16x4_t s[2];
292      s[0] = vld1q_u8_x4((uint8_t*)&c[0]);
293      s[1] = vld1q_u8_x4((uint8_t*)&c[8]);
294      gift_64_vec_sliced_bits_pack(s);
295
296      uint8x16x4_t round_keys[ROUNDS_GIFT_64][2];
297      gift_64_vec_sliced_generate_round_keys(round_keys, key);
298
299      for (int round = ROUNDS_GIFT_64 - 1; round >= 0; round--) {
300          s[0].val[0] = veorq_u8(s[0].val[0], round_keys[round][0].val[0]);
301          s[0].val[1] = veorq_u8(s[0].val[1], round_keys[round][0].val[1]);
302          s[0].val[2] = veorq_u8(s[0].val[2], round_keys[round][0].val[2]);
303          s[0].val[3] = veorq_u8(s[0].val[3], round_keys[round][0].val[3]);
304          s[1].val[0] = veorq_u8(s[1].val[0], round_keys[round][1].val[0]);
305          s[1].val[1] = veorq_u8(s[1].val[1], round_keys[round][1].val[1]);
306          s[1].val[2] = veorq_u8(s[1].val[2], round_keys[round][1].val[2]);
307          s[1].val[3] = veorq_u8(s[1].val[3], round_keys[round][1].val[3]);
308
309          gift_64_vec_sliced_permute_inv(s);
310          gift_64_vec_sliced_subcells_inv(s);
311      }
312
313      gift_64_vec_sliced_bits_pack(s);
314      vst1q_u8_x4((uint8_t*)&m[0], s[0]);
315      vst1q_u8_x4((uint8_t*)&m[8], s[1]);
316  }
```

Bibliography

[1] Subhadeep Banik et al. "GIFT: A Small Present". In: Aug. 2017, pp. 321–345. ISBN: 978-3-319-66786-7. DOI: `10.1007/978-3-319-66787-4_16`.

[2] Ryad Benadjila et al. "Implementing Lightweight Block Ciphers on x86 Architectures". In: Selected Areas in Cryptography – SAC 2013. Ed. by Tanja Lange, Kristin Lauter, and Petr Lisoněk. Berlin, Heidelberg: Springer Berlin Heidelberg, 2014, pp. 324–351. ISBN: 978-3-662-43414-7.