

Efficient Implementation Strategies for Block Ciphers on ARMv8

Bachelorarbeit

Bastian Engel

April 11, 2023

Abstract

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Declaration

I hereby declare that ...

Contents

1	Introduction	5
1.1	Notation	5
1.2	Block ciphers	5
1.2.1	GIFT	6
1.2.2	Camellia	8
1.3	The ARMv8 platform	9
1.3.1	General architecture	9
1.3.2	NEON	10
1.3.3	NEON Intrinsics	12
2	Implementation strategies	14
2.1	Strategies for GIFT	14
2.1.1	Table-based	14
2.1.2	Using <code>vperm</code>	15
2.1.3	Bitslicing	16
2.2	Strategies for Camellia	21
3	Implementation	22
3.1	Pipelining	22
3.2	GIFT	23
3.2.1	Table-based	23
3.2.2	Using <code>vperm</code>	23
3.2.3	Bitslicing	24
4	Evaluation	27
4.1	Benchmarks	27
4.1.1	GIFT	27
4.1.2	Camellia	27
A	C implementations	31
A.1	Implementations for SPN	31

A.1.1	Table-based	31
A.1.2	Using <code>vperm</code>	34
A.1.3	Bitslicing	39
B	Lorem dolor	46

Chapter 1

Introduction

1.1 Notation

1.2 Block ciphers

Securing communication channels between different parties has been a long-term subject of study for cryptographers and engineers which is essential to our modern world to cope with ever-increasing amounts of devices producing and sharing data. The main way to facilitate high-throughput, confidential communications nowadays is through the use of symmetric cryptography in which two parties share a common secret, called a key, which allows them to encrypt, share and subsequently decrypt messages to achieve confidentiality against third parties. Ciphers can be divided into two categories; block ciphers, which always encrypt fixed-sized messages called blocks, and stream ciphers, which continuously provide encryption for an arbitrarily long, constant stream of data.

A block cipher can be defined as a bijection between the input block (the message) and the output block (the ciphertext). For any block cipher with block size n , we denote the key-dependent encryption and decryption functions as $E_K, D_K : \mathbb{F}_2^n \rightarrow \mathbb{F}_2^n$. The simplest way to characterize this bijection is through a lookup table which yields the highest possible performance as each block can be encrypted by one simple lookup depending on the key and the message. This is not practical though due to most ciphers working with block and key sizes $n, |K| \geq 64$. For a block cipher with $n = 64, |K| = 128$, a space of $2^{64}2^{128}64 = 2^{198}$ is necessary. Considering modern consumer hard disks being able to store data in the order of 2^{40} , it is easy to see that a lookup table is wholly impractical. We therefore describe block ciphers algorithmically which opens up possibilities for different tradeoffs and security concerns.

1.2.1 GIFT

GIFT[1], first presented in the *CHES 2017* cryptographic hardware and embedded systems conference, is a lightweight block cipher based on a previous design called **PRESENT**, developed in 2007. Its goal is to offer maximum security while being extremely light on resources. Modern battery-powered devices like RFID tags or low-latency operations like on-the-fly disc encryption present strong hardware and power constraints. **GIFT** aims to be a simple, low-energy cipher suited for these kinds of applications.

GIFT comes in two variants; **GIFT-64** working with 64-bit blocks and **GIFT-128** working with 128-bit blocks. In both cases, the key is 128 bits long. The design is a very simple, round-based substitution-permutation network (SPN). One round consists in a sequential application of the confusion layer by means of 4-bit S-boxes and subsequent diffusion through bit permutation. After the bit permutation, a round key is added to the cipher state and the single round is complete. **GIFT-64** uses 28 rounds while **GIFT-128** uses 40 rounds.




Figure 1.1: Two rounds of **GIFT-64**

Substitution layer

The input of **GIFT** is split into 4-bit nibbles which are then fed into 16 S-boxes for **GIFT-64** and 32 S-boxes for **GIFT-128**. The S-box $S : \mathbb{F}_2^4 \rightarrow \mathbb{F}_2^4$ is defined as follows:

x	0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f
$S(x)$	1	a	4	c	6	f	3	9	2	d	b	7	5	0	8	e

Permutation layer

The permutation P works on individual bits and maps bit b_i to $b_{P(i)}$. The different permutations for GIFT-64 and GIFT-128 can be expressed by:

$$P_{64}(i) = 4 \left\lfloor \frac{i}{16} \right\rfloor + 16 \left(\left(3 \left\lfloor \frac{i \bmod 16}{4} \right\rfloor + (i \bmod 4) \right) \bmod 4 \right) + (i \bmod 4)$$

$$P_{128}(i) = 4 \left\lfloor \frac{i}{16} \right\rfloor + 32 \left(\left(3 \left\lfloor \frac{i \bmod 16}{4} \right\rfloor + (i \bmod 4) \right) \bmod 4 \right) + (i \bmod 4)$$

Round key addition

The last step of each round consists in XORing a round key R_i to the cipher state. The new cipher state x_{i+1} after each full round is therefore given by

$$x_{i+1} = P(S(x_i)) \oplus R_i$$

Round key extraction and key schedule

Round key extraction differs for GIFT-64 and GIFT-128. Let $K_{(128)} = k_7 || k_6 || \dots || k_0$ denote the 128-bit key state.

GIFT-64 . We extract two words $U_{(16)} || V_{(16)} = k_1 || k_0$ from the key state. These are then added to round key $R_{(64)}$: $R_{4i+1} \leftarrow U_i, R_{4i} \leftarrow V_i$.

GIFT-128 . We extract two words $U_{(32)} || V_{(32)} = k_5 || k_4 || k_1 || k_0$ from the key state. These are then added to round key $R_{(128)}$: $R_{4i+2} \leftarrow U_i, R_{4i+1} \leftarrow V_i$.

In both cases, we additionally XOR a round constant $C_{(6)}$ to bit positions $n - 1, 23, 19, 15, 11, 7, 3$. The round constants are generated using a 6-bit affine linear-feedback shift register and have the following values:

Rounds	Constants
1 - 16	01, 03, 07, 0F, 1F, 3E, 3D, 3B, 37, 2F, 1E, 3C, 39, 33, 27, 0E
17 - 32	1D, 3A, 35, 2B, 16, 2C, 18, 30, 21, 02, 05, 0B, 17, 2E, 1C, 38
33 - 48	31, 23, 06, 0D, 1B, 36, 2D, 1A, 34, 29, 12, 24, 08, 11, 22, 04

The key state is then updated by individually rotating k_1 and k_0 and rotating the new state 32 bits to the right:

$$k_7 || k_6 || \dots || k_1 || k_0 \leftarrow k_1 \ggg 2 || k_0 \ggg 12 || k_7 || k_6 || \dots || k_3 || k_2$$

1.2.2 Camellia

Camellia[2] is a block cipher jointly developed by NTT and Mitsubishi Electric Corporation and first published in 2001. Following AES specifications, it is able to encrypt 128-bit blocks using either 128-, 192- or 256-bit keys and claims to possess similar performance and security levels as the AES finalists.

Encryption overview

The encryption process has an 18-round Feistel structure for 128-bit keys and a 24-round Feistel structure for 192/256-bit key and employs key whitening to increase security. First, subkeys $kw_{t(64)}(t = 0, 1, 2, 3)$, $k_{u(64)}(u = 0, 1, \dots, (17|23))$ and $kl_{v(64)}(v = 0, 1, 2, 3)$ are generated from the master key. Then, pre-whitening keys are applied to the plaintext $m_{(128)} = L_{(64)} || R_{(64)}$:

$$(L || R) \leftarrow (L || R) \oplus (kw_0 || kw_1)$$

The next steps differ for 128-bit and 192/256-bit keys in the number of rounds:

Table 1.1: Camellia encryption

Round r	128-bit keys	192/256-bit keys
0-4	$(L R) \leftarrow FE(L, R, k_r)$	$(L R) \leftarrow FE(L, R, k_r)$
5	$(L R) \leftarrow FE(L, R)$	$(L R) \leftarrow FE(L, R)$
	$(L R) \leftarrow FLL(L, R, kl_0, kl_1)$	$(L R) \leftarrow FLL(L, R, kl_0, kl_1)$
6-10	$(L R) \leftarrow FE(L, R, k_r)$	$(L R) \leftarrow FE(L, R, k_r)$
11	$(L R) \leftarrow FE(L, R)$	$(L R) \leftarrow FE(L, R)$
	$(L R) \leftarrow FLL(L, R, kl_2, kl_3)$	$(L R) \leftarrow FLL(L, R, kl_2, kl_3)$
12-16	$(L R) \leftarrow FE(L, R, k_r)$	$(L R) \leftarrow FE(L, R, k_r)$
17		$(L R) \leftarrow FE(L, R)$
		$(L R) \leftarrow FLL(L, R, kl_4, kl_5)$
18-23		$(L R) \leftarrow FE(L, R, k_r)$

Finally, R and L are concatenated and XORed with the post-whitening keys to obtain the cipher text $c_{(128)}$:

$$c = (R||L) \oplus (kw_2)||kw_3)$$

Components

We will give an overview of the main functional components of Camellia.

Feistel round function FE

$$\begin{aligned} FE : (\mathbb{F}_2^{64})^3 &\rightarrow (\mathbb{F}_2^{64})^2 \\ (L_{(64)}, R_{(64)}, k_{(64)}) &\mapsto (R \oplus F(L, k), L) \end{aligned}$$

FL layer function FLL

$$\begin{aligned} FLL : (\mathbb{F}_2^{64})^4 &\rightarrow (\mathbb{F}_2^{64})^2 \\ (L_{(64)}, R_{(64)}, k_{0(64)}, k_{1(64)}) &\mapsto (FL(L, k_0), FL^{-1}(R, k_1)) \end{aligned}$$

1.3 The ARMv8 platform

With small devices and embedded processors becoming ever more ubiquitous and essential in areas like consumer electronics or industrial and IoT applications, the need for low-power, high-performance microprocessors has increased steadily. With more than 250 billion chips shipped, semiconductors designed by ARM power 95% of mobile devices and have found a great many applications due to their high performance and low power consumption[3]. The ODROID-N2+[4] development board we are using is based on the big.LITTLE architecture and is powered by a quad-core ARM Cortex-A73 processor and a weaker dual-core ARM Cortex-A53 for power efficiency. Both these processors are part of the eight generation of ARM designs known as ARMv8[5].

ARMv8 defines three architecture profiles for different use cases as well as dynamic execution states with corresponding instruction sets. This work will focus on the A profile running in the AArch64 state utilizing the A64 instruction set with NEON and crypto extensions.

1.3.1 General architecture

ARMv8 is a RISC architecture employing simple data processing instructions operating only on registers as well as dedicated load/store instructions to transfer

Table 1.2: ARMv8 profiles

Profile	Description
Application (A)	Traditional use with virtual memory and privilege level support
Real-time (R)	Real-time, low-latency, deterministic embedded systems
Microcontroller (M)	Very low-power, fast-interrupt embedded systems

Table 1.3: ARMv8 execution states

Execution state	Usage	Instruction sets
AArch32	32-bit compatibility	A32/T32
AArch64	64-bit	A64

data from register to memory and back. This enables faster execution of individual instructions, a simpler pipeline design, predictable instruction timings and fewer addressing modes.

The A64 instruction set defines 31 64-bit general-purpose registers **X0–X30** which can also be accessed as 32-bit registers **W0–W30**. Values are loaded from and stored to memory using **LDR/STR**. Data processing instructions generally use explicit output registers instead of overwriting the first input register.

Table 1.4: A64 addressing modes

Addressing mode	Example	Description
Base register	LDR W0, [X1]	W0 = *(X1);
Offset	LDR W0, [X1, #12]	W0 = *(X1 + 12);
Pre-indexing	LDR W0, [X1, #12]!	X1 += 12; W0 = *(X1);
Post-indexing	LDR W0, [X1], #12	W0 = *(X1); X1 += 12;

1.3.2 NEON

ARMv8 supports single-instruction, multiple-data (SIMD) processing. These systems allow the programmer to store multiple pieces of data in a vector and work on them in parallel to speed up calculations. The A64 instruction set defines two possible SIMD implementations:

1. Advanced SIMD, known as NEON

2. Scalable Vector Extension (SVE)

We will take a look at NEON as this is the type of vector processing supported by the Cortex-A73 processor.

The register file of the NEON unit is made up of 32 quad-word (128-bit) registers **V0–V31**, each extending the standard 64-bit floating-point registers **D0–31**. These registers are divided into equally sized lanes on which vector instructions operate. Figure 1.2 shows valid ways to interpret the register **V0**.




Figure 1.2: Divisions of the V0 register

NEON instructions interpret their operands' layouts (i.e. lane count and width) through the use of suffixes such as **.4S** or **.8H**. Adding eight 16-bit halfwords stored in **V1** and **V2** can be done as follows:

ADD V0.8H, V1.8H, V2.8H




Figure 1.3: Addition of two vector registers

1.3.3 NEON Intrinsics

The header file `<arm_neon.h>` provides ARM-specific data and function definitions including vector data types and C functions for working with these vectors. These functions are known as NEON intrinsics [6] and give the programmer a high-level interface to most NEON instructions. Major advantages of this approach include the ease of development as the compiler takes over register allocation and load/store operations as well as performance benefits through compiler optimizations.

Standard vector data types have the format `uintn \times m_t` with lane width n in bits and lane count m . Array types of the format `uintn \times m \times c_t`, $c \in \{2, 3, 4\}$ are also defined which are used in operations requiring multiple parameters like `TBL` or pairwise load/stores. Intrinsics include the operation name and lane data format as well as an optional `q` suffix to indicate operation on a 128-bit register. Multiplying eight pairs of 16-bit numbers `a, b` for example can be done via the following:

```
uint16 $\times$ 8_t result = vmulq_u16(a, b);
```

In this case, the compiler allocates vector registers for `a`, `b` and `result` and assembles the intrinsic to `MUL Vr.8H, Va.8H, Vb.8H`. Necessary loads and stores for the result and parameters are also handled automatically. Of special interest to us are the following intrinsics, each existing in different variants with different lane widths and also array types:

Table 1.5: Common NEON intrinsics

Intrinsic		Summary
<code>uint64_t</code>	<code>vgetq_lane_u64(void)</code>	Extract a single lane
<code>void</code>	<code>vsetq_lane_u64(uint64_t)</code>	Insert a single lane
<code>uint64x2_t</code>	<code>vdupq_n_u64(uint64_t)</code>	Initialize all lanes to same value
<code>void</code>	<code>vst1q_u64(uint64_t*, uint64x2_t)</code>	Store from register to memory
<code>uint64x2_t</code>	<code>vld1q_u64(uint64_t*, uint64x2_t)</code>	Load from memory to register
<code>uint8x16_t</code>	<code>veorq_u8(uint8x16_t, uint8x16_t)</code>	bitwise XOR
<code>uint8x16_t</code>	<code>vandq_u8(uint8x16_t, uint8x16_t)</code>	bitwise AND
<code>uint8x16_t</code>	<code>vorrq_u8(uint8x16_t, uint8x16_t)</code>	bitwise OR
<code>uint8x16_t</code>	<code>vmvnq_u8(uint8x16_t)</code>	bitwise NOT
<code>uint8x16_t</code>	<code>vqtbl2q_u8(uint8x16_t, uint8x16_t)</code>	permutation (TBL)

Chapter 2

Implementation strategies

Due to the structural differences of SPN- and Feistel network-based ciphers, we shall analyze these two separately.

2.1 Strategies for **GIFT**

Three implementation strategies for substitution-permutation networks are introduced by [7]:

- Table-based implementations
- **vperm** implementations
- Bitslice implementations

2.1.1 Table-based

Table-driven programming is a simple way to increase performance of operations by tabulating the results, therefore requiring only a single memory access to acquire the result. This approach is obviously limited to manageable table sizes, so while tabulating a function like the AES S-box $S_{AES} : \mathbb{F}_2^8 \rightarrow \mathbb{F}_2^8$ requires only 2^{11} space, tabulating the **GIFT** permutation layer $P_{GIFT} : \mathbb{F}_2^{64} \rightarrow \mathbb{F}_2^{64}$ would require 2^{70} space, which is totally unfeasible.

A common approach is to tabulate the output of each S-box, including the diffusion layer, and then XORing the results together. Let n denote the internal cipher state size and s the size of a single S-box in bits. For each S-box $S_i, i \in \{0, \dots, \frac{n}{s}\}$, we can construct a mapping $T_i : \mathbb{F}_2^s \rightarrow \mathbb{F}_2^n$ representing substitution with subsequent permutation of that single S-box. The cipher state before round key addition is then given by $\bigoplus_{i=0}^{\frac{n}{s}-1} T_i(m_i)$ for each s -bit message chunk m_i . This

approach requires space of $\frac{n}{s}|\mathbb{F}_2^s|n = \frac{n^2 2^s}{s}$ bits, which, for **GIFT-64**, results in a manageable size of $\frac{64^2 2^4}{4} = 2^{14}$ bits which equals 16 KiB.

Constructing the tables

For **GIFT-64**, table construction is relatively straightforward and can be done as follows:

Listing 2.1: Table construction algorithm

```

1  tables <- [[]]
2  for sbox_index from 0 to 15 do
3      for sbox_input from 0 to 15 do
4          output <- sbox(sbox_input)
5          output <- permute(output << (4 * sbox_index))
6          tables[sbox_index][sbox_input] <- output

```

Implementing this algorithm gives us the following table representing the first and second S-box.

x	$T_0(x)$	$T_1(x)$...
$0x0$	$0x1$	$0x1000000000000$...
$0x1$	$0x8000000020000$	$0x800000002$...
$0x2$	$0x400000000$	$0x40000$...
$0x3$	$0x8000400000000$	$0x800040000$...
$0x4$	$0x400020000$	$0x40002$...
$0x5$	$0x8000400020001$	$0x1000800040002$...
$0x6$	$0x20001$	$0x1000000000002$...
$0x7$	$0x80000000000001$	$0x1000800000000$...
$0x8$	$0x20000$	$0x2$...
$0x9$	$0x8000400000001$	$0x1000800040000$...
$0xa$	$0x8000000020001$	$0x1000800000002$...
$0xb$	$0x400020001$	$0x1000000040002$...
$0xc$	$0x400000001$	$0x1000000040000$...
$0xd$	$0x0$	$0x0$...
$0xe$	$0x80000000000000$	$0x800000000$...
$0xf$	$0x8000400020000$	$0x800040002$...

The tables for **GIFT-128** can be generated in a similar way by looping through all 32 S-boxes instead of 16 on line 3.

2.1.2 Using **vperm**

The plenitude of different processing instructions introduced by NEON1.3.2 allow flexible ways to further speed up algorithms having reached their optimizational

limit on non-SIMD platforms. **vperm**, a general term standing for *vector permute*, is a common instruction on SIMD machines. Called **TBL** on NEON, it is used for parallel table lookups and arbitrary permutations. It takes two inputs to perform a lanewise lookup:

1. A register with lookup values
2. One or more registers containing data

S-box lookup

This instruction can be used to implement S-box lookup of all 16 S-boxes in a single instruction. We do this by packing our 64-bit cipher state $s = s_{15}||s_{14}||\dots||s_0$ into a vector register V_0 . Because we can only operate on whole bytes, we put each 4-bit S-box into an 8-bit lane which neatly fits into the 128-bit registers. We then put the S-box itself into register V_1 which will be used as the data register for the table lookup.

The confusion layer can now be performed through one **TBL** instruction:

TBL V0.16B, V1.16B, V0.16B




Figure 2.1: Performing the S-Box lookup in parallel

2.1.3 Bitslicing

Bitslicing refers to the technique of splitting up n bits into m slices to achieve a more efficient representation to operate on. The structure of **GIFT** naturally offers possibilities for bitslicing. We split the cipher state bits $b_{63}b_{62}\dots b_0$ into four slices $S_i, i \in \{0, 1, 2, 3\}$ such that the i -th slice contains all i -th bits of the individual S-boxes. This is equivalent to transposing the bit matrix.

$$S = \begin{bmatrix} S_0 \\ S_1 \\ S_2 \\ S_3 \end{bmatrix} = \begin{bmatrix} b_{60}b_{56}b_{52} \dots b_0 \\ b_{61}b_{57}b_{53} \dots b_1 \\ b_{62}b_{58}b_{54} \dots b_2 \\ b_{63}b_{59}b_{55} \dots b_3 \end{bmatrix}$$

Parallel S-Boxes

This representation offers multiple advantages. We first note that computation of the S-box can be executed in parallel, similar to the `vperm` technique above. This can be done by finding an algorithmic way to apply the S-box which has already been proposed by the original **GIFT** authors:

$$\begin{aligned} S_1 &\leftarrow S_1 \oplus (S_0 \wedge S_2) \\ t &\leftarrow S_0 \oplus (S_1 \wedge S_3) \\ S_2 &\leftarrow S_2 \oplus (t \vee S_1) \\ S_0 &\leftarrow S_3 \oplus S_2 \\ S_1 &\leftarrow S_1 \oplus S_0 \\ S_0 &\leftarrow \neg S_0 \\ S_2 &\leftarrow S_2 \oplus (t \wedge S_1) \\ S_3 &\leftarrow t \end{aligned}$$

This is very efficient as it only requires six XOR-, three AND and one OR operation.

An important property of the permutation is the fact that bits always stay in their slice. This means we can decompose the permutation P into four permutations $P_i, i \in \{0, 1, 2, 3\}$ and apply these permutations separately to each slice. One possible way to implement a permutation P_i in software is to mask off all bits individually, shift them to their correct position and OR them together:

$$P_i(S_i) = \bigvee_{k=0}^{15} (S_i \wedge m_i) \ll s_i$$

This approach requires 47 operations, meaning all four permutations require over 150 operations which would present a major bottleneck to the round function. We can improve on this by working on multiple message blocks at once and using the aforementioned `vperm` instruction to implement the bit shuffling. We then need only four instructions for the complete diffusion layer.

Using **vperm** for slice permutation

We cannot use the **TBL** instruction directly as we need to shuffle individual bits, but the smallest data we can operate on are bytes. We therefore encrypt $8n$ messages at once which allows us to create bitwise groupings. These messages are put into $4m$ registers with register R_{4i} containing S_0 , register R_{4i+1} containing S_1 and so forth. With block size BS and register size RS , the following must hold:

$$8n \cdot BS = 4m \cdot RS$$

In the case of **GIFT-64** with $BS = 64$ and ARM NEON with $RS = 128$, we get

$$8n \cdot 64 = 4m \cdot 128 \Leftrightarrow n = m$$

$n = m = 1$ would be a valid choice which yields eight messages divided into four registers. We choose $n = m = 2$ so we can directly utilize the algorithm for bit packing presented by the original GIFT authors, although it is simple to adapt this algorithm to only four registers and eight messages by adjusting the **SWAPMOVE** shift and mask values.

Packing the data into bitslice format

Let a, b, \dots, p be sixteen messages of length 64 with subscripts denoting individual bits. We first put these messages into eight SIMD registers V_0, V_1, \dots, V_7 :

$$\begin{aligned} V_0 &= b||a & V_4 &= j||i \\ V_1 &= d||c & V_5 &= l||k \\ V_2 &= f||e & V_6 &= n||m \\ V_3 &= h||g & V_7 &= p||o \end{aligned}$$

We then use the **SWAPMOVE** technique to bring the data into bitslice format. This operation operates on two registers A, B using mask M and shift value N . It swaps bits in A masked by $(M \ll N)$ with bits in B masked by M in using only three XOR-, one AND- and two shift operations.

$$\begin{aligned} &\text{SWAPMOVE}(A, B, M, N) : \\ &T = ((A \gg N) \oplus B) \wedge M \\ &B = B \oplus T \\ &A = A \oplus (T \ll N) \end{aligned}$$

One caveat of this approach is the fact that NEON registers cannot be shifted in their entirety due to the fact bits are not able to cross lanes. This leads to the problem of being able to shift at most two lanes of 64 bits at once. We thus need to implement the $\text{shr}(V, n)$ and $\text{shl}(V, n)$ operations on our own. This can be done by first extracting the 64-bit lanes a, b out of $V = b||a$, shifting the lanes individually and finally shifting and ORing the crossing bits back into the other lane.

$$\begin{aligned} \text{shl}(V, n) : \\ a, b &= V[0], V[1] \\ c &= (a \gg (64 - n)) \\ V[0] &= (a \ll n) \\ V[1] &= (b \ll n) \vee c \end{aligned}$$

The following operations group all i -th bits of the messages a, c, \dots, o into bytes and puts these into the lower half of the registers $V_{i \bmod 8}$. The same is done for messages b, d, \dots, p , only differing in that the bytes are put into the upper half of the registers.

SWAPMOVE($V_0, V_1, 0x5555 \dots 55, 1$) SWAPMOVE($V_4, V_5, 0x5555 \dots 55, 1$)
 SWAPMOVE($V_2, V_3, 0x5555 \dots 55, 1$) SWAPMOVE($V_6, V_7, 0x5555 \dots 55, 1$)
 SWAPMOVE($V_0, V_2, 0x3333 \dots 33, 2$) SWAPMOVE($V_4, V_6, 0x3333 \dots 33, 2$)
 SWAPMOVE($V_1, V_3, 0x3333 \dots 33, 2$) SWAPMOVE($V_5, V_7, 0x3333 \dots 33, 2$)
 SWAPMOVE($V_0, V_4, 0xf0f0 \dots 0f, 4$) SWAPMOVE($V_1, V_5, 0xf0f0 \dots 0f, 4$)
 SWAPMOVE($V_2, V_6, 0xf0f0 \dots 0f, 4$) SWAPMOVE($V_3, V_7, 0xf0f0 \dots 0f, 4$)

With $Ax = o_x m_x k_x j_x g_x e_x c_x a_x$ and $Bx = p_x n_x l_x i_x h_x f_x d_x b_x$ denoting byte groups, our data now has the following permutation-friendly format:

n	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
V_0	$B56$	$B48$	$B40$	$B32$	$B24$	$B16$	$B8$	$B0$	$A56$	$A48$	$A40$	$A32$	$A24$	$A16$	$A8$	$A0$
V_1	$B57$	$B49$	$B41$	$B33$	$B25$	$B17$	$B9$	$B1$	$A57$	$A49$	$A41$	$A33$	$A25$	$A17$	$A9$	$A1$
V_2	$B58$	$B50$	$B42$	$B34$	$B26$	$B18$	$B10$	$B2$	$A58$	$A50$	$A42$	$A34$	$A26$	$A18$	$A10$	$A2$
V_3	$B59$	$B51$	$B43$	$B35$	$B27$	$B19$	$B11$	$B3$	$A59$	$A51$	$A43$	$A35$	$A27$	$A19$	$A11$	$A3$
V_4	$B60$	$B52$	$B44$	$B36$	$B28$	$B20$	$B12$	$B4$	$A60$	$A52$	$A44$	$A36$	$A28$	$A20$	$A12$	$A4$
V_5	$B61$	$B53$	$B45$	$B37$	$B29$	$B21$	$B13$	$B5$	$A61$	$A53$	$A45$	$A37$	$A29$	$A21$	$A13$	$A5$
V_6	$B62$	$B54$	$B46$	$B38$	$B30$	$B22$	$B14$	$B6$	$A62$	$A54$	$A46$	$A38$	$A30$	$A22$	$A14$	$A6$
V_7	$B63$	$B55$	$B47$	$B39$	$B31$	$B23$	$B15$	$B7$	$A63$	$A55$	$A47$	$A39$	$A31$	$A23$	$A15$	$A7$

Although this would already work, we prefer to have only bits of the same messages in each register - otherwise the permutation would need to operate on two

source registers with the added requirement of storing the pre-permutation values for the first four registers, slowing down the round function through superfluous load/stores. This transformation is trivial by use of **TBL** with two data source operands. The final data format we operate on is as follows:

n	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
V_0	A60	A56	A52	A48	A44	A40	A36	A32	A28	A24	A20	A16	A12	A8	A4	A0
V_1	A61	A57	A53	A49	A45	A41	A37	A33	A29	A25	A21	A17	A13	A9	A5	A1
V_2	A62	A58	A54	A50	A46	A42	A38	A34	A30	A26	A22	A18	A14	A10	A6	A2
V_3	A63	A59	A55	A51	A47	A43	A39	A35	A31	A27	A23	A19	A15	A11	A7	A3
V_4	B60	B56	B52	B48	B44	B40	B36	B32	B28	B24	B20	B16	B12	B8	B4	B0
V_5	B61	B57	B53	B49	B45	B41	B37	B33	B29	B25	B21	B17	B13	B9	B5	B1
V_6	B62	B58	B54	B50	B46	B42	B38	B34	B30	B26	B22	B18	B14	B10	B6	B2
V_7	B63	B59	B55	B51	B47	B43	B39	B35	B31	B27	B23	B19	B15	B11	B7	B3

We can now create permutation tables using the specification of the individual slice permutations P_i which are then applied to V_i and V_{i+4} respectively:

j	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
$P_0(j)$	0	12	8	4	1	13	9	5	2	14	10	6	3	15	11	7
$P_1(j)$	4	0	12	8	5	1	13	9	6	2	14	10	7	3	15	11
$P_2(j)$	8	4	0	12	9	5	1	13	10	6	2	14	11	7	3	15
$P_3(j)$	12	8	4	0	13	9	5	1	14	10	6	2	15	11	7	3

One thing to take note of is the original permutation values only show where a given byte should land, not which byte belongs to a certain position - i.e. for P_0 , byte 1 should land in position 12, but the byte belonging to position 1 is byte 4. Because **TBL** works in the latter way, we have to do some trivial rearrangements.

Assuming the correct permutation values are put into registers V_8, V_9, V_{10}, V_{11} , this now allows us to compute the permutation layer for all 16 blocks in only eight permutation instructions.

TBL V0, V0, V8	TBL V1, V1, V9
TBL V4, V4, V8	TBL V5, V5, V9
TBL V2, V2, V10	TBL V3, V3, V11
TBL V6, V6, V10	TBL V7, V7, V11

Round key function

In contrast to packing and unpacking of data which is only done once in the beginning and end, a round key is derived for every round, so the round key derivation function needs to be as fast as possible. A simple but naive approach

for one round would be to generate a single round key, copy it 15 times and pack the resulting registers similar to how we proceed with the messages. Due to the cost of packing the messages, this is prohibitively expensive. Because we know where each byte group ends up after packing, we can directly XOR the round key bits to the correct position. Extending these bits to bytes can then be done simply by repeatedly shifting and ORing the registers together.

2.2 Strategies for Camellia

Chapter 3

Implementation

Implementations in the C programming language for the presented strategies can be found in Appendix A. Although directly writing Assembler code could result in a small performance benefit, this generally increases the work necessary by an order of magnitude for only limited results. Instruction-level optimization and in particular register allocation is left to the compiler. Relying on the compiler mandates a closer study of the generated, optimized assembler. All source files were compiled using clang version 15.0.7 and optimization level O3.

3.1 Pipelining

Understanding certain choices requires an understanding of the Cortex-A73 instruction pipeline[8]. Being a superscalar processor, it is able to execute more than one instruction per clock cycle by dispatching instructions to different execution units working in parallel.

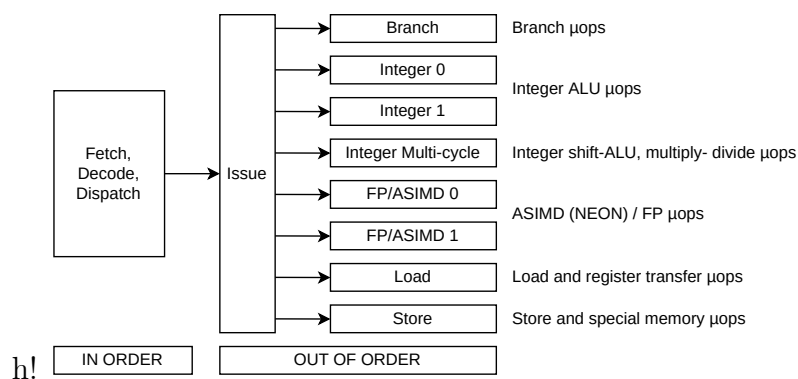


Figure 3.1: High-level overview of the Cortex-A72 instruction pipeline

The processor might for example store a calculation result, load a necessary value from memory and execute two SIMD operations at once, all in the same clock cycle. Modern compilers take advantage of this fact by reordering instructions such that all pipeline execution units stay as busy as possible and do not stall while having to wait for new instructions to be dispatched. A more thorough analysis will be presented for the bitsliced strategy of **GIFT**, but all implementations are heavily reliant on function inlining, instruction reordering and loop unrolling.

3.2 GIFT

3.2.1 Table-based

This is the simplest strategy to implement. Indeed, its biggest advantage lies in its portability to other platforms without relying on specific features or extensions. The cipher state is stored as a 64-bit word and one round consists in extracting the 4-bit S-boxes, looking up table values, collecting these in an accumulator and finally adding the round key.

```

1 uint64_t gift_64_table_subperm(const uint64_t cipher_state)
2 {
3     uint64_t new_cipher_state = 0;
4
5     for (size_t i = 0; i < 16; i++) {
6         int nibble = (cipher_state >> (i * 4)) & 0xf;
7         new_cipher_state ^= tables[i][nibble];
8     }
9
10    return new_cipher_state;
11 }
```

```

1 uint64_t gift_64_table_encrypt(const uint64_t m,
2                               const uint64_t rks[restrict ROUNDS_GIFT_64])
3 {
4     uint64_t c = m;
5
6     // round loop
7     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
8         c = gift_64_table_subperm(c);
9         c ^= rks[round];
10    }
11
12    return c;
13 }
```

3.2.2 Using **vperm**

Implementation of the substitution layer requires the use of a single vector intrinsic. This mandates the packing of data into a vector register which in turn is

disadvantageous to the permutation layer as we need to extract single bits. Packing and unpacking is nothing more than filling 8-bit vector lanes with 4-bit S-boxes and vice versa.

```

1 uint8x16_t gift_64_vec_sbox_subcells(const uint8x16_t cipher_state)
2 {
3     return vqtbl1q_u8(sbox_vec, cipher_state);
4 }

1 uint8x16_t gift_64_vec_sbox_permute(const uint8x16_t cipher_state)
2 {
3     // collect individual bits into 64-bit register
4     uint64_t new_cipher_state = 0UL;
5     uint64_t boxes[2];
6     vst1q_u64(boxes, cipher_state);
7
8     for (size_t box = 0; box < 16; box++) {
9         for (size_t i = 0; i < 4; i++) {
10             const int bit = (boxes[box / 8] >> ((box % 8) * 8 + i)) & 0x1;
11             new_cipher_state |= (uint64_t)bit << perm_64[box * 4 + i];
12         }
13     }
14
15     return gift_64_vec_sbox_bits_pack(new_cipher_state);
16 }

1 uint64_t gift_64_vec_sbox_encrypt(const uint64_t m,
2                                 const uint8x16_t rks[restrict ROUNDS_GIFT_64])
3 {
4     // pack into vector register
5     uint8x16_t c = gift_64_vec_sbox_bits_pack(m);
6
7     // round loop
8     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
9         c = gift_64_vec_sbox_subcells(c);
10        c = gift_64_vec_sbox_permute(c);
11        c = veorq_u8(c, rks[round]);
12    }
13
14    // unpack
15    return gift_64_vec_sbox_bits_unpack(c);
16 }

```

3.2.3 Bitslicing

We will examine the round function in closer detail and compare the source code with the generated assembly.

Listing 3.1: Round function src

```

1  for (int round = 0; round < ROUNDS_GIFT_64; round++) {
2      gift_64_vec_sliced_subcells(s);
3      gift_64_vec_sliced_permute(s);
4
5      // round key addition
6      s[0].val[0] = veorq_u8(s[0].val[0], rks[round][0].val[0]);
7      s[0].val[1] = veorq_u8(s[0].val[1], rks[round][0].val[1]);
8      s[0].val[2] = veorq_u8(s[0].val[2], rks[round][0].val[2]);
9      s[0].val[3] = veorq_u8(s[0].val[3], rks[round][0].val[3]);
10     s[1].val[0] = veorq_u8(s[1].val[0], rks[round][1].val[0]);
11     s[1].val[1] = veorq_u8(s[1].val[1], rks[round][1].val[1]);
12     s[1].val[2] = veorq_u8(s[1].val[2], rks[round][1].val[2]);
13     s[1].val[3] = veorq_u8(s[1].val[3], rks[round][1].val[3]);
14 }

```

Listing 3.2: Round function asm

1	and	v20.16b, v17.16b, v6.16b	24	tbl	v18.16b, {v6.16b}, v2.16b
2	add	x9, x20, x8	25	tbl	v19.16b, {v21.16b}, v3.16b
3	eor	v20.16b, v20.16b, v16.16b	26	tbl	v21.16b, {v4.16b}, v2.16b
4	add	x8, x8, #0x80	27	ldp	q6, q4, [x9, #-112]
5	and	v21.16b, v20.16b, v19.16b	28	mvn	v16.16b, v16.16b
6	cmp	x8, #0xe70	29	tbl	v16.16b, {v16.16b}, v0.16b
7	eor	v21.16b, v21.16b, v6.16b	30	tbl	v17.16b, {v17.16b}, v1.16b
8	orr	v6.16b, v6.16b, v16.16b	31	mvn	v5.16b, v5.16b
9	eor	v6.16b, v6.16b, v17.16b	32	tbl	v5.16b, {v5.16b}, v0.16b
10	eor	v16.16b, v19.16b, v6.16b	33	ldp	q22, q23, [x9, #-80]
11	eor	v17.16b, v16.16b, v20.16b	34	eor	v6.16b, v6.16b, v16.16b
12	and	v19.16b, v21.16b, v17.16b	35	eor	v16.16b, v4.16b, v17.16b
13	eor	v6.16b, v19.16b, v6.16b	36	tbl	v7.16b, {v7.16b}, v1.16b
14	and	v19.16b, v7.16b, v4.16b	37	eor	v17.16b, v22.16b, v18.16b
15	eor	v19.16b, v19.16b, v5.16b	38	tbl	v20.16b, {v20.16b}, v3.16b
16	and	v20.16b, v19.16b, v18.16b	39	ldp	q4, q18, [x9, #-48]
17	eor	v20.16b, v20.16b, v4.16b	40	eor	v19.16b, v23.16b, v19.16b
18	orr	v4.16b, v4.16b, v5.16b	41	eor	v4.16b, v4.16b, v5.16b
19	eor	v4.16b, v4.16b, v7.16b	42	ldp	q22, q23, [x9, #-16]
20	eor	v5.16b, v18.16b, v4.16b	43	eor	v5.16b, v18.16b, v7.16b
21	eor	v7.16b, v5.16b, v19.16b	44	eor	v7.16b, v22.16b, v21.16b
22	and	v18.16b, v20.16b, v7.16b	45	eor	v18.16b, v23.16b, v20.16b
23	eor	v4.16b, v18.16b, v4.16b	46	b.ne	197e0

Without needing to dive too deep into details, it is obvious to see the compiler has inlined the two function calls to **subcells** and **permute** with lines 1 to 23 originating from the **subcells** and lines 24-46 from the **permute** and round key addition functions. It has chosen not to unroll the loop, but has moved the loop counter increment as well as the condition check in between the **subcells** instructions to line 4 and 6. In addition, the loop counter serves a second purpose as an offset register and is therefore incremented by 0x80=128 instead of just 1.

Permutation (**tbl**) and round key addition (**eor**) instructions are interleaved. The compiler recognizes data dependencies and can therefore proceed with round key addition immediately after a slice has been permuted without needing to wait for all permutations to finish. This is only logical considering the inner workings of the instruction pipeline: by interleaving NEON with regular logic and load instruc-

tions, the execution units are filled more evenly and pipeline stalls are prevented which speeds up computation.

Round keys are loaded from memory a few instructions before they are needed; assuming all round keys are stored in the L1 cache, loading a floating-point/vector register takes 5 cycles. After the load has been dispatched to the load execution unit in line 27, the processor happily continues chugging away at the instruction stream by issuing `tbl` and `mvn` μ ops to other execution units.

These kinds of optimizations are pervasive when programming using higher level languages like C and modern-day compilers more often than not outperform handcrafted assembly.

Chapter 4

Evaluation

In this chapter, we will evaluate the strategies through performance measurements and discuss advantages, disadvantages and possible use cases.

4.1 Benchmarks

Performance measurements were taken for each strategy as well as for naive reference implementations and are presented through latency *lat* in cycles per byte (c/B) as well as constant throughput *thr* in MiB/s of the entire encryption strategy. Round key derivation is measured separately. Measurements of all individual components like packing or permuting have to be viewed as upper bounds due to the aforementioned inlining, instruction reordering and pipelining (TODO REF HERE) taking place in the actual encryption function.

The AArch64 defines system registers in addition to general-purpose registers which are used for system configuration and monitoring. One of these registers is the performance monitor cycle count register **PMCCNTR** which counts processor clock cycles. Access from userspace is disabled by default and can be activated through a custom Linux kernel module by setting **PMUSERENR.EN** to 1. To minimize interference and because the cycle count register is core-local, we isolate and utilize one Cortex-A53 and Cortex-A73 core from the rest of the system for exclusive benchmarking purposes respectively by use of the **isolcpus** kernel command line parameter and **taskset** command utility.

4.1.1 GIFT

4.1.2 Camellia

Table 4.1: Benchmarks for GIFT

Strategy	Component	Cortex-A53		Cortex-A73	
		<i>lat</i> (c/B)	<i>thr</i> (MiB/s)	<i>lat</i> (c/B)	<i>thr</i> (MiB/s)
Naive GIFT-64	round_keys	223.54	1.22	190.34	2.33
	encrypt	1367.66		830.49	
	subcells	4.64		3.13	
	permute	36.68		21.09	
Naive GIFT-128	round_keys	182.91	0.51	167.16	0.79
	encrypt	3532.52		2615.89	
	subcells	6.91		2.73	
	permute	82.68		61.27	
Table-driven	round_keys	223.81	13.59	190.11	16.09
	encrypt	122.11		119.62	
	subperm	5.15		4.47	
vperm S-box	round_keys	308.12	2.53	270.80	3.93
	encrypt	658.95		492.29	
	subcells	1.63		1.13	
	permute	22.71		15.24	
	pack	7.51		2.94	
	unpack	8.76		5.89	
Bitsliced	round_keys	10.13	111.89	7.66	154.21
	encrypt	16.22		13.51	
	subcells	0.41		0.18	
	permute	0.39		0.09	
	pack	1.48		1.01	
	unpack	1.49		1.01	

Table 4.2: Benchmarks for Camellia

Strategy	Component	Cortex-A53		Cortex-A73	
		<i>lat</i> (c/B)	<i>thr</i> (MiB/s)	<i>lat</i> (c/B)	<i>thr</i> (MiB/s)
Naive GIFT-64	round_keys	223.54	1.22	190.34	2.33
	encrypt	1367.66		830.49	
	subcells	4.64		3.13	
	permute	36.68		21.09	
Naive GIFT-128	round_keys	182.91	0.51	167.16	0.79
	encrypt	3532.52		2615.89	
	subcells	6.91		2.73	
	permute	82.68		61.27	
Table-driven	round_keys	223.81	13.59	190.11	16.09
	encrypt	122.11		119.62	
	subperm	5.15		4.47	
vperm S-box	round_keys	308.12	2.53	270.80	3.93
	encrypt	658.95		492.29	
	subcells	1.63		1.13	
	permute	22.71		15.24	
	pack	7.51		2.94	
	unpack	8.76		5.89	
Bitsliced	round_keys	10.13	111.89	7.66	154.21
	encrypt	16.22		13.51	
	subcells	0.41		0.18	
	permute	0.39		0.09	
	pack	1.48		1.01	
	unpack	1.49		1.01	

Acknowledgements

I want to thank ...

Appendix A

C implementations

A.1 Implementations for SPN

A.1.1 Table-based

Listing A.1: gift_table.h

```
1  #ifndef GIFT_TABLE_H
2  #define GIFT_TABLE_H
3
4  #include <stdint.h>
5
6  #define ROUNDS_GIFT_64 28
7
8  void gift_64_table_generate_round_keys(uint64_t rks[restrict ROUNDS_GIFT_64],
9                                         const uint64_t key[restrict 2]);
10
11 uint64_t gift_64_table_subperm(const uint64_t cipher_state);
12
13 // can only encrypt using table technique!
14 uint64_t gift_64_table_encrypt(const uint64_t m,
15                               const uint64_t rks[restrict ROUNDS_GIFT_64]);
16
17 #endif
```

Listing A.2: gift_table.c

```
1  #include "gift_table.h"
2
3  #include <stdlib.h>
4  #include <string.h>
5
6  static const int round_const[] = {
7      // rounds 0-15
8      0x01, 0x03, 0x07, 0x0F, 0x1F, 0x3E, 0x3D, 0x3B,
9      0x37, 0x2F, 0x1E, 0x3C, 0x39, 0x33, 0x27, 0x0E,
10     // rounds 16-31
11     0x1D, 0x3A, 0x35, 0x2B, 0x16, 0x2C, 0x18, 0x30,
12     0x21, 0x02, 0x05, 0x0B, 0x17, 0x2E, 0x1C, 0x38,
13     // rounds 32-47
```



```

14         0x31, 0x23, 0x06, 0x0D, 0x1B, 0x36, 0x2D, 0x1A,
15         0x34, 0x29, 0x12, 0x24, 0x08, 0x11, 0x22, 0x04
16     };
17
18     static const uint64_t tables[16][16] = {
19         { 0x0000000000000001UL, 0x00080000000020000UL, 0x0000000040000000UL, 0
          x0008000400000000UL, 0x0000000040020000UL, 0x0008000400020001UL, 0
          x00000000000020001UL, 0x00080000000000001UL, 0x00000000000020000UL, 0
          x00080004000000001UL, 0x00080000000020001UL, 0x00000000400020001UL, 0
          x00000000400000001UL, 0x0000000000000000UL, 0x0008000000000000UL, 0
          x0008000400020000UL },
20         { 0x0001000000000000UL, 0x00000000800000002UL, 0x0000000000004000UL, 0
          x00000000800040000UL, 0x00000000000040002UL, 0x0001000800040002UL, 0
          x00010000000000002UL, 0x0001000800000000UL, 0x00000000000000002UL, 0
          x0001000800040000UL, 0x00010008000000002UL, 0x00010000000040002UL, 0
          x0001000000040000UL, 0x0000000000000000UL, 0x0000000080000000UL, 0
          x00000000800040002UL },
21         { 0x0000000100000000UL, 0x0002000000008000UL, 0x0000000000000004UL, 0
          x0000000000008000UL, 0x0002000000000004UL, 0x0002000100080004UL, 0
          x0002000100000000UL, 0x0000000010008000UL, 0x0002000000000000UL, 0
          x0000000100080004UL, 0x0002000100080000UL, 0x0002000100000004UL, 0
          x00000000100000004UL, 0x0000000000000000UL, 0x0000000000008000UL, 0
          x00020000000080004UL },
22         { 0x0000000000010000UL, 0x00000000200000008UL, 0x0004000000000000UL, 0
          x0004000000000008UL, 0x0004000200000000UL, 0x0004000200010008UL, 0
          x00000000200010000UL, 0x00000000000010008UL, 0x0000000020000000UL, 0
          x0004000000010008UL, 0x0000000020010008UL, 0x0004000200010000UL, 0
          x0004000000010000UL, 0x0000000000000000UL, 0x0000000000000008UL, 0
          x0004000200000008UL },
23         { 0x0000000000000001UL, 0x0080000000200000UL, 0x0000000400000000UL, 0
          x0080004000000000UL, 0x0000004000200000UL, 0x0080004000200010UL, 0
          x0000000000200010UL, 0x0080000000000010UL, 0x000000000020000UL, 0
          x0080004000000001UL, 0x0080000000200010UL, 0x00000004000200010UL, 0
          x0000004000000001UL, 0x0000000000000000UL, 0x0080000000000000UL, 0
          x0080004000200000UL },
24         { 0x0010000000000000UL, 0x0000008000000020UL, 0x0000000000400000UL, 0
          x0000008000400000UL, 0x0000000000400020UL, 0x0010008000400020UL, 0
          x0010000000000020UL, 0x0010008000000000UL, 0x0000000000000020UL, 0
          x0010008000400000UL, 0x0010008000000020UL, 0x0010000000400020UL, 0
          x0010000000400000UL, 0x0000000000000000UL, 0x0000008000000000UL, 0
          x0000008000400020UL },
25         { 0x0000001000000000UL, 0x0020000000800000UL, 0x000000000000040UL, 0
          x0000000000800040UL, 0x002000000000040UL, 0x0020001000800040UL, 0
          x0020001000000000UL, 0x0000001000800000UL, 0x0020000000000000UL, 0
          x0000001000800040UL, 0x0020001000800000UL, 0x0020001000000040UL, 0
          x0000001000000040UL, 0x0000000000000000UL, 0x0000000000800000UL, 0
          x0020000000800040UL },
26         { 0x0000000000100000UL, 0x0000002000000080UL, 0x0040000000000000UL, 0
          x0040000000000080UL, 0x0040002000000000UL, 0x0040002000100080UL, 0
          x0000002000100000UL, 0x0000000000100080UL, 0x0000002000000000UL, 0
          x0040000000100080UL, 0x0000002000100080UL, 0x0040002000100000UL, 0
          x0040000000100000UL, 0x0000000000000000UL, 0x00000000000080UL, 0
          x0040002000000080UL },
27         { 0x0000000000000001UL, 0x0800000000200000UL, 0x0000004000000000UL, 0
          x0800040000000000UL, 0x0000004000200000UL, 0x0800040002000100UL, 0
          x0000000002000100UL, 0x0800000000000010UL, 0x0000000002000000UL, 0
          x0800040000000010UL, 0x0800000002000100UL, 0x0000004000200010UL, 0
          x0000040000000010UL, 0x0000000000000000UL, 0x0800000000000000UL, 0
          x0800040002000000UL },
28         { 0x0100000000000000UL, 0x0000008000000020UL, 0x0000000000400000UL, 0
          x0000008000400000UL, 0x0000000000400020UL, 0x0100080004000200UL, 0
          x0100000000000020UL, 0x0100080000000000UL, 0x000000000000200UL, 0

```

```

    x0100080004000000UL, 0x0100080000000200UL, 0x0100000004000200UL, 0
    x0100000004000000UL, 0x0000000000000000UL, 0x0000080000000000UL, 0
    x0000080004000200UL },
29 { 0x0000010000000000UL, 0x0200000008000000UL, 0x0000000000000400UL, 0
    x0000000008000400UL, 0x0200000000000400UL, 0x0200010008000400UL, 0
    x0200010000000000UL, 0x0000010008000000UL, 0x0200000000000000UL, 0
    x0000010008000400UL, 0x0200010008000000UL, 0x0200010000000400UL, 0
    x0000010000000400UL, 0x0000000000000000UL, 0x0000000008000000UL, 0
    x0200000008000400UL },
30 { 0x0000000001000000UL, 0x0000020000000080UL, 0x0400000000000000UL, 0
    x0400000000000800UL, 0x0400020000000000UL, 0x0400020001000800UL, 0
    x0000020001000000UL, 0x0000000001000800UL, 0x0000020000000000UL, 0
    x0400000001000800UL, 0x0000020001000800UL, 0x0400020001000000UL, 0
    x0400000001000000UL, 0x0000000000000000UL, 0x0000000000000800UL, 0
    x0400020000000800UL },
31 { 0x000000000001000UL, 0x8000000020000000UL, 0x0000400000000000UL, 0
    x8000400000000000UL, 0x0000400020000000UL, 0x8000400020001000UL, 0
    x0000000020001000UL, 0x800000000001000UL, 0x0000000020000000UL, 0
    x8000400000001000UL, 0x8000000020001000UL, 0x0000400020001000UL, 0
    x0000400000001000UL, 0x0000000000000000UL, 0x8000000000000000UL, 0
    x8000400020000000UL },
32 { 0x1000000000000000UL, 0x0000800000002000UL, 0x0000000040000000UL, 0
    x0000800040000000UL, 0x0000000040002000UL, 0x1000800040002000UL, 0
    x1000000000002000UL, 0x1000800000000000UL, 0x0000000000002000UL, 0
    x1000800040000000UL, 0x1000800000002000UL, 0x1000000040002000UL, 0
    x1000000040000000UL, 0x0000000000000000UL, 0x0000800000000000UL, 0
    x0000800040002000UL },
33 { 0x0000100000000000UL, 0x2000000080000000UL, 0x0000000000004000UL, 0
    x0000000080004000UL, 0x2000000000004000UL, 0x2000100080004000UL, 0
    x2000100000000000UL, 0x0000100080000000UL, 0x2000000000000000UL, 0
    x0000100080004000UL, 0x2000100080000000UL, 0x2000100000004000UL, 0
    x0000100000004000UL, 0x0000000000000000UL, 0x0000000080000000UL, 0
    x2000000080004000UL },
34 { 0x0000000010000000UL, 0x0000200000008000UL, 0x4000000000000000UL, 0
    x400000000008000UL, 0x4000200000000000UL, 0x4000200010008000UL, 0
    x0000200010000000UL, 0x0000000010008000UL, 0x0000200000000000UL, 0
    x4000000010008000UL, 0x0000200010008000UL, 0x4000200010000000UL, 0
    x4000000010000000UL, 0x0000000000000000UL, 0x000000000008000UL, 0
    x400020000008000UL }
35 };
36
37 void gift_64_table_generate_round_keys(uint64_t rks[restrict ROUNDS_GIFT_64],
38                                     const uint64_t key[restrict 2])
39 {
40     uint64_t key_state[] = {key[0], key[1]};
41     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
42         int v = (key_state[0] >> 0) & 0xffff;
43         int u = (key_state[0] >> 16) & 0xffff;
44
45         // add round key (RK=U|V)
46         rks[round] = 0UL;
47         for (size_t i = 0; i < 16; i++) {
48             int key_bit_v = (v >> i) & 0x1;
49             int key_bit_u = (u >> i) & 0x1;
50             rks[round] ^= (uint64_t)key_bit_v << (i * 4 + 0);
51             rks[round] ^= (uint64_t)key_bit_u << (i * 4 + 1);
52         }
53
54         // add single bit
55         rks[round] ^= 1UL << 63;
56
57         // add round constants

```

```

58         rks[round] ^= ((round_const[round] >> 0) & 0x1) << 3;
59         rks[round] ^= ((round_const[round] >> 1) & 0x1) << 7;
60         rks[round] ^= ((round_const[round] >> 2) & 0x1) << 11;
61         rks[round] ^= ((round_const[round] >> 3) & 0x1) << 15;
62         rks[round] ^= ((round_const[round] >> 4) & 0x1) << 19;
63         rks[round] ^= ((round_const[round] >> 5) & 0x1) << 23;
64
65         // update key state
66         int k0 = (key_state[0] >> 0) & 0xffffUL;
67         int k1 = (key_state[0] >> 16) & 0xffffUL;
68         k0 = (k0 >> 12) | ((k0 & 0xfff) << 4);
69         k1 = (k1 >> 2) | ((k1 & 0x3) << 14);
70         key_state[0] >>= 32;
71         key_state[0] |= (key_state[1] & 0xffffffffUL) << 32;
72         key_state[1] >>= 32;
73         key_state[1] |= ((uint64_t)k0 << 32) | ((uint64_t)k1 << 48);
74     }
75 }
76
77 uint64_t gift_64_table_subperm(const uint64_t cipher_state)
78 {
79     uint64_t new_cipher_state = 0;
80
81     for (size_t i = 0; i < 16; i++) {
82         int nibble = (cipher_state >> (i * 4)) & 0xf;
83         new_cipher_state ^= tables[i][nibble];
84     }
85
86     return new_cipher_state;
87 }
88
89 uint64_t gift_64_table_encrypt(const uint64_t m,
90                               const uint64_t rks[restrict ROUNDS_GIFT_64])
91 {
92     uint64_t c = m;
93
94     // round loop
95     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
96         c = gift_64_table_subperm(c);
97         c ^= rks[round];
98     }
99
100     return c;
101 }

```

A.1.2 Using vperm

Listing A.3: gift_vec_sbox.h

```

1  #ifndef GIFT_VEC_SBOX_H
2  #define GIFT_VEC_SBOX_H
3
4  #include <stdint.h>
5  #include <arm_neon.h>
6
7  #define ROUNDS_GIFT_64 28
8
9  // expose for benchmarking

```

```

10 uint8x16_t gift_64_vec_sbox_bits_pack(const uint64_t a);
11 uint64_t gift_64_vec_sbox_bits_unpack(const uint8x16_t a);
12 uint8x16_t gift_64_vec_sbox_subcells(const uint8x16_t cipher_state);
13 uint8x16_t gift_64_vec_sbox_subcells_inv(const uint8x16_t cipher_state);
14 uint8x16_t gift_64_vec_sbox_permute(const uint8x16_t cipher_state);
15 uint8x16_t gift_64_vec_sbox_permute_inv(const uint8x16_t cipher_state);
16 void gift_64_vec_sbox_generate_round_keys(uint8x16_t rks[ROUNDS_GIFT_64],
17                                           const uint64_t key[restrict 2]);
18
19 // construct tables
20 void gift_64_vec_sbox_init(void);
21
22 uint64_t gift_64_vec_sbox_encrypt(const uint64_t m,
23                                  const uint8x16_t rks[restrict ROUNDS_GIFT_64]);
24 uint64_t gift_64_vec_sbox_decrypt(const uint64_t c,
25                                   const uint8x16_t rks[restrict ROUNDS_GIFT_64]);
26
27 #endif

```

Listing A.4: gift_vec_sbox.c

```

1  #include "gift_vec_sbox.h"
2
3  #include <stdint.h>
4  #include <arm_neon.h>
5  #include <string.h>
6
7  static uint64_t sbox_vec_u64[2] = {
8      0x09030f060c040a01UL, 0x0e080005070b0d02UL
9  };
10
11 static uint64_t sbox_vec_inv_u64[2] = {
12     0x0b040c020608000dUL, 0x050f09030a01070eUL
13 };
14
15 static uint8x16_t sbox_vec;
16 static uint8x16_t sbox_vec_inv;
17
18 // split S-box bits into vector lanes
19 uint8x16_t gift_64_vec_sbox_bits_pack(const uint64_t a)
20 {
21     uint8x16_t v;
22     v = vsetq_lane_u64(
23         (uint64_t)((a >> 4 * 0) & 0xf) << 8 * 0 |
24         (uint64_t)((a >> 4 * 1) & 0xf) << 8 * 1 |
25         (uint64_t)((a >> 4 * 2) & 0xf) << 8 * 2 |
26         (uint64_t)((a >> 4 * 3) & 0xf) << 8 * 3 |
27         (uint64_t)((a >> 4 * 4) & 0xf) << 8 * 4 |
28         (uint64_t)((a >> 4 * 5) & 0xf) << 8 * 5 |
29         (uint64_t)((a >> 4 * 6) & 0xf) << 8 * 6 |
30         (uint64_t)((a >> 4 * 7) & 0xf) << 8 * 7, v, 0);
31
32     v = vsetq_lane_u64(
33         (uint64_t)((a >> 4 * 8) & 0xf) << 8 * 0 |
34         (uint64_t)((a >> 4 * 9) & 0xf) << 8 * 1 |
35         (uint64_t)((a >> 4 * 10) & 0xf) << 8 * 2 |
36         (uint64_t)((a >> 4 * 11) & 0xf) << 8 * 3 |
37         (uint64_t)((a >> 4 * 12) & 0xf) << 8 * 4 |
38         (uint64_t)((a >> 4 * 13) & 0xf) << 8 * 5 |
39         (uint64_t)((a >> 4 * 14) & 0xf) << 8 * 6 |
40         (uint64_t)((a >> 4 * 15) & 0xf) << 8 * 7, v, 1);

```

```

41     return v;
42 }
43
44 // merge S-box bits into single uint64_t
45 uint64_t gift_64_vec_sbox_bits_unpack(const uint8x16_t v)
46 {
47     uint64_t a = 0UL;
48     a |= (uint64_t)vgetq_lane_u8(v, 0) << 4 * 0;
49     a |= (uint64_t)vgetq_lane_u8(v, 1) << 4 * 1;
50     a |= (uint64_t)vgetq_lane_u8(v, 2) << 4 * 2;
51     a |= (uint64_t)vgetq_lane_u8(v, 3) << 4 * 3;
52     a |= (uint64_t)vgetq_lane_u8(v, 4) << 4 * 4;
53     a |= (uint64_t)vgetq_lane_u8(v, 5) << 4 * 5;
54     a |= (uint64_t)vgetq_lane_u8(v, 6) << 4 * 6;
55     a |= (uint64_t)vgetq_lane_u8(v, 7) << 4 * 7;
56     a |= (uint64_t)vgetq_lane_u8(v, 8) << 4 * 8;
57     a |= (uint64_t)vgetq_lane_u8(v, 9) << 4 * 9;
58     a |= (uint64_t)vgetq_lane_u8(v, 10) << 4 * 10;
59     a |= (uint64_t)vgetq_lane_u8(v, 11) << 4 * 11;
60     a |= (uint64_t)vgetq_lane_u8(v, 12) << 4 * 12;
61     a |= (uint64_t)vgetq_lane_u8(v, 13) << 4 * 13;
62     a |= (uint64_t)vgetq_lane_u8(v, 14) << 4 * 14;
63     a |= (uint64_t)vgetq_lane_u8(v, 15) << 4 * 15;
64
65     return a;
66 }
67
68 static const size_t perm_64[] = {
69     0, 17, 34, 51, 48, 1, 18, 35, 32, 49, 2, 19, 16, 33, 50, 3,
70     4, 21, 38, 55, 52, 5, 22, 39, 36, 53, 6, 23, 20, 37, 54, 7,
71     8, 25, 42, 59, 56, 9, 26, 43, 40, 57, 10, 27, 24, 41, 58, 11,
72     12, 29, 46, 63, 60, 13, 30, 47, 44, 61, 14, 31, 28, 45, 62, 15
73 };
74
75 static const size_t perm_64_inv[] = {
76     0, 5, 10, 15, 16, 21, 26, 31, 32, 37, 42, 47, 48, 53, 58, 63,
77     12, 1, 6, 11, 28, 17, 22, 27, 44, 33, 38, 43, 60, 49, 54, 59,
78     8, 13, 2, 7, 24, 29, 18, 23, 40, 45, 34, 39, 56, 61, 50, 55,
79     4, 9, 14, 3, 20, 25, 30, 19, 36, 41, 46, 35, 52, 57, 62, 51
80 };
81
82 static const int round_const[] = {
83     // rounds 0-15
84     0x01, 0x03, 0x07, 0x0F, 0x1F, 0x3E, 0x3D, 0x3B, 0x37, 0x2F, 0x1E, 0x3C, 0x39, 0
85     x33, 0x27, 0x0E,
86     // rounds 16-31
87     0x1D, 0x3A, 0x35, 0x2B, 0x16, 0x2C, 0x18, 0x30, 0x21, 0x02, 0x05, 0x0B, 0x17, 0
88     x2E, 0x1C, 0x38,
89     // rounds 32-47
90     0x31, 0x23, 0x06, 0x0D, 0x1B, 0x36, 0x2D, 0x1A, 0x34, 0x29, 0x12, 0x24, 0x08, 0
91     x11, 0x22, 0x04
92 };
93
94 uint8x16_t gift_64_vec_sbox_subcells(const uint8x16_t cipher_state)
95 {
96     return vqtbl1q_u8(sbox_vec, cipher_state);
97 }
98
99 uint8x16_t gift_64_vec_sbox_subcells_inv(const uint8x16_t cipher_state)
100 {
101     return vqtbl1q_u8(sbox_vec_inv, cipher_state);
102 }

```

```

100 }
101
102 uint8x16_t gift_64_vec_sbox_permute(const uint8x16_t cipher_state)
103 {
104     // collect individual bits into 64-bit register
105     uint64_t new_cipher_state = 0UL;
106     uint64_t boxes[2];
107     vst1q_u64(boxes, cipher_state);
108
109     for (size_t box = 0; box < 16; box++) {
110         for (size_t i = 0; i < 4; i++) {
111             const int bit = (boxes[box / 8] >> ((box % 8) * 8 + i)) & 0x1;
112             new_cipher_state |= (uint64_t)bit << perm_64[box * 4 + i];
113         }
114     }
115
116     return gift_64_vec_sbox_bits_pack(new_cipher_state);
117 }
118
119 uint8x16_t gift_64_vec_sbox_permute_inv(const uint8x16_t cipher_state)
120 {
121     // collect into 64-bit register (faster)
122     uint64_t new_cipher_state = 0;
123     uint64_t boxes[2];
124     vst1q_u64(boxes, cipher_state);
125
126     // S-box 0-7
127     for (size_t box = 0; box < 8; box++) {
128         for (size_t i = 0; i < 4; i++) {
129             const int bit = (boxes[0] >> (box * 8 + i)) & 0x1;
130             new_cipher_state |= (uint64_t)bit << perm_64_inv[box * 4 + i];
131         }
132     }
133
134     // S-box 8-15
135     for (size_t box = 0; box < 8; box++) {
136         for (size_t i = 0; i < 4; i++) {
137             const int bit = (boxes[1] >> (box * 8 + i)) & 0x1;
138             new_cipher_state |= (uint64_t)bit << perm_64_inv[(box + 8) * 4
139                 + i];
140         }
141     }
142
143     return gift_64_vec_sbox_bits_pack(new_cipher_state);
144 }
145
146 void gift_64_vec_sbox_generate_round_keys(uint8x16_t rks[ROUNDS_GIFT_64],
147     const uint64_t key[2])
148 {
149     uint64_t key_state[] = {key[0], key[1]};
150     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
151         const int v = (key_state[0] >> 0) & 0xffff;
152         const int u = (key_state[0] >> 16) & 0xffff;
153
154         // add round key (RK=U||V)
155         uint64_t round_key = 0UL;
156         for (size_t i = 0; i < 16; i++) {
157             const int key_bit_v = (v >> i) & 0x1;
158             const int key_bit_u = (u >> i) & 0x1;
159             round_key ^= (uint64_t)key_bit_v << (i * 4 + 0);
160             round_key ^= (uint64_t)key_bit_u << (i * 4 + 1);
161         }
162     }
163 }

```

```

161
162         // add single bit
163         round_key ^= 1UL << 63;
164
165         // add round constants
166         round_key ^= ((round_const[round] >> 0) & 0x1) << 3;
167         round_key ^= ((round_const[round] >> 1) & 0x1) << 7;
168         round_key ^= ((round_const[round] >> 2) & 0x1) << 11;
169         round_key ^= ((round_const[round] >> 3) & 0x1) << 15;
170         round_key ^= ((round_const[round] >> 4) & 0x1) << 19;
171         round_key ^= ((round_const[round] >> 5) & 0x1) << 23;
172
173         // pack into vector register
174         rks[round] = gift_64_vec_sbox_bits_pack(round_key);
175
176         // update key state
177         int k0 = (key_state[0] >> 0) & 0xffffUL;
178         int k1 = (key_state[0] >> 16) & 0xffffUL;
179         k0 = (k0 >> 12) | ((k0 & 0xfff) << 4);
180         k1 = (k1 >> 2) | ((k1 & 0x3) << 14);
181         key_state[0] >>= 32;
182         key_state[0] |= (key_state[1] & 0xffffffffUL) << 32;
183         key_state[1] >>= 32;
184         key_state[1] |= ((uint64_t)k0 << 32) | ((uint64_t)k1 << 48);
185     }
186 }
187
188 void gift_64_vec_sbox_init(void)
189 {
190     // construct sbox_vec
191     sbox_vec = vld1q_u64(sbox_vec_u64);
192
193     // construct sbox_vec_inv
194     sbox_vec_inv = vld1q_u64(sbox_vec_inv_u64);
195 }
196
197 uint64_t gift_64_vec_sbox_encrypt(const uint64_t m,
198                                 const uint8x16_t rks[restrict ROUNDS_GIFT_64])
199 {
200     // pack into vector register
201     uint8x16_t c = gift_64_vec_sbox_bits_pack(m);
202
203     // round loop
204     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
205         c = gift_64_vec_sbox_subcells(c);
206         c = gift_64_vec_sbox_permute(c);
207         c = veorq_u8(c, rks[round]);
208     }
209
210     // unpack
211     return gift_64_vec_sbox_bits_unpack(c);
212 }
213
214 uint64_t gift_64_vec_sbox_decrypt(const uint64_t c,
215                                 const uint8x16_t rks[restrict ROUNDS_GIFT_64])
216 {
217     // pack into vector register
218     uint8x16_t m = gift_64_vec_sbox_bits_pack(c);
219
220     // round loop (in reverse)
221     for (int round = ROUNDS_GIFT_64 - 1; round >= 0; round--) {
222         m = veorq_u8(m, rks[round]);

```

```

223         m = gift_64_vec_sbox_permute_inv(m);
224         m = gift_64_vec_sbox_subcells_inv(m);
225     }
226
227     // unpack
228     return gift_64_vec_sbox_bits_unpack(m);
229 }

```

A.1.3 Bitslicing

Listing A.5: gift_vec_sliced.h

```

1  #ifndef GIFT_VEC_SLICED_H
2  #define GIFT_VEC_SLICED_H
3
4  #include <stdint.h>
5  #include <arm_neon.h>
6
7  #define ROUNDS_GIFT_64 28
8
9  // expose for benchmarking
10 uint8x16_t shl(const uint8x16_t v, const int n);
11 uint8x16_t shr(const uint8x16_t v, const int n);
12 void gift_64_vec_sliced_swapmove(uint8x16_t *restrict a, uint8x16_t *restrict b,
13                                 const uint8x16_t m, const int n);
14 void gift_64_vec_sliced_bits_pack(uint8x16x4_t m[restrict 2]);
15 void gift_64_vec_sliced_bits_unpack(uint8x16x4_t m[restrict 2]);
16
17 void gift_64_vec_sliced_subcells(uint8x16x4_t cipher_state[restrict 2]);
18 void gift_64_vec_sliced_subcells_inv(uint8x16x4_t cipher_state[restrict 2]);
19 void gift_64_vec_sliced_permute(uint8x16x4_t cipher_state[restrict 2]);
20 void gift_64_vec_sliced_permute_inv(uint8x16x4_t cipher_state[2]);
21 void gift_64_vec_sliced_generate_round_keys(uint8x16x4_t rks[restrict ROUNDS_GIFT_64
22                                             ][2],
23                                             const uint64_t key[restrict 2]);
24
25 void gift_64_vec_sliced_init(void);
26
27 void gift_64_vec_sliced_encrypt(uint64_t c[restrict 16],
28                                const uint64_t m[restrict 16],
29                                const uint8x16x4_t rks[restrict ROUNDS_GIFT_64][2]);
30 void gift_64_vec_sliced_decrypt(uint64_t m[restrict 16],
31                                 const uint64_t c[restrict 16],
32                                 const uint8x16x4_t rks[restrict ROUNDS_GIFT_64][2]);
33
34 #endif

```

Listing A.6: gift_vec_sliced.c

```

1  #include "gift_vec_sliced.h"
2
3  #include <stddef.h>
4  #include <stdio.h>
5  #include <string.h>
6
7  static uint64_t pack_shf_u64[4] = {
8      0x1303120211011000UL, 0x1707160615051404UL, // S0/S1/S2/S3
9      0x1b0b1a0a19091808UL, 0x1f0f1e0e1d0d1c0cUL, // S4/S5/S6/S7

```



```

10 };
11
12 static uint64_t pack_shf_inv_u64[4] = {
13     0x0e0c0a0806040200UL, 0x1e1c1a1816141210UL, // S0/S1/S2/S3
14     0x0f0d0b0907050301UL, 0x1f1d1b1917151311UL, // S4/S5/S6/S7
15 };
16
17 static uint64_t perm_u64[8] = {
18     0x0f0b07030c080400UL, 0x0d0905010e0a0602UL, // S0/S4
19     0x0c0804000d090501UL, 0x0e0a06020f0b0703UL, // S1/S5
20     0x0d0905010e0a0602UL, 0x0f0b07030c080400UL, // S2/S6
21     0x0e0a06020f0b0703UL, 0x0c0804000d090501UL // S3/S7
22 };
23
24
25 static uint64_t perm_inv_u64[8] = {
26     0x05090d0104080c00UL, 0x070b0f03060a0e02UL, // S0/S4
27     0x090d0105080c0004UL, 0x0b0f03070a0e0206UL, // S1/S5
28     0x0d0105090c000408UL, 0x0f03070b0e02060aUL, // S2/S6
29     0x0105090d0004080cUL, 0x03070b0f02060a0eUL // S3/S7
30 };
31
32 static uint8x16x2_t pack_shf;
33 static uint8x16x2_t pack_shf_inv;
34 static uint8x16x4_t perm;
35 static uint8x16x4_t perm_inv;
36
37 static uint8x16_t pack_mask_0;
38 static uint8x16_t pack_mask_1;
39 static uint8x16_t pack_mask_2;
40
41 static const int round_const[] = {
42     // rounds 0-15
43     0x01, 0x03, 0x07, 0x0F, 0x1F, 0x3E, 0x3D, 0x3B, 0x37, 0x2F, 0x1E, 0x3C, 0x39, 0
44     x33, 0x27, 0x0E,
45     // rounds 16-31
46     0x1D, 0x3A, 0x35, 0x2B, 0x16, 0x2C, 0x18, 0x30, 0x21, 0x02, 0x05, 0x0B, 0x17, 0
47     x2E, 0x1C, 0x38,
48     // rounds 32-47
49     0x31, 0x23, 0x06, 0x0D, 0x1B, 0x36, 0x2D, 0x1A, 0x34, 0x29, 0x12, 0x24, 0x08, 0
50     x11, 0x22, 0x04
51 };
52
53 uint8x16_t shl(const uint8x16_t v, const int n)
54 {
55     uint64_t l[2];
56     vst1q_u64(l, v);
57     l[1] = (l[1] << n) | (l[0] >> (64 - n));
58     l[0] <<= n;
59     return vld1q_u64(l);
60 }
61
62 uint8x16_t shr(const uint8x16_t v, const int n)
63 {
64     uint64_t l[2];
65     vst1q_u64(l, v);
66     l[0] = l[0] >> n | (l[1] << (64 - n));
67     l[1] >>= n;
68     return vld1q_u64(l);
69 }
70
71 void gift_64_vec_sliced_swapmove(uint8x16_t *restrict a, uint8x16_t *restrict b,

```

```

69         const uint8x16_t m, const int n)
70     {
71
72         const uint8x16_t t = vandq_u8(veorq_u8(shr(*a, n), *b), m);
73         *b = veorq_u8(*b, t);
74         *a = veorq_u8(*a, shl(t, n));
75     }
76
77 void gift_64_vec_sliced_bits_pack(uint8x16x4_t m[restrict 2])
78 {
79     // take care not to shift mask bits out of the register
80     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[1], pack_mask_0, 1);
81     gift_64_vec_sliced_swapmove(&m[0].val[2], &m[0].val[3], pack_mask_0, 1);
82     gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[1], pack_mask_0, 1);
83     gift_64_vec_sliced_swapmove(&m[1].val[2], &m[1].val[3], pack_mask_0, 1);
84
85     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[2], pack_mask_1, 2);
86     gift_64_vec_sliced_swapmove(&m[0].val[1], &m[0].val[3], pack_mask_1, 2);
87     gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[2], pack_mask_1, 2);
88     gift_64_vec_sliced_swapmove(&m[1].val[1], &m[1].val[3], pack_mask_1, 2);
89
90     // make bytes (a0 b0 c0 d0 a4 b4 c4 d4 -> a0 b0 c0 d0 e0 f0 g0 h0)
91     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[1].val[0], pack_mask_2, 4);
92     gift_64_vec_sliced_swapmove(&m[0].val[2], &m[1].val[2], pack_mask_2, 4);
93     gift_64_vec_sliced_swapmove(&m[0].val[1], &m[1].val[1], pack_mask_2, 4);
94     gift_64_vec_sliced_swapmove(&m[0].val[3], &m[1].val[3], pack_mask_2, 4);
95
96     // same plaintext slice bits into same register (so we only have to do
97     // what we are doing here once instead of every round)
98     const uint8x16x2_t pairs[4] = {
99         { .val = { m[0].val[0], m[1].val[0] } },
100        { .val = { m[0].val[1], m[1].val[1] } },
101        { .val = { m[0].val[2], m[1].val[2] } },
102        { .val = { m[0].val[3], m[1].val[3] } },
103    };
104
105    m[0].val[0] = vqtbl2q_u8(pairs[0], pack_shf.val[0]);
106    m[0].val[1] = vqtbl2q_u8(pairs[1], pack_shf.val[0]);
107    m[0].val[2] = vqtbl2q_u8(pairs[2], pack_shf.val[0]);
108    m[0].val[3] = vqtbl2q_u8(pairs[3], pack_shf.val[0]);
109
110    m[1].val[0] = vqtbl2q_u8(pairs[0], pack_shf.val[1]);
111    m[1].val[1] = vqtbl2q_u8(pairs[1], pack_shf.val[1]);
112    m[1].val[2] = vqtbl2q_u8(pairs[2], pack_shf.val[1]);
113    m[1].val[3] = vqtbl2q_u8(pairs[3], pack_shf.val[1]);
114 }
115
116 void gift_64_vec_sliced_bits_unpack(uint8x16x4_t m[restrict 2])
117 {
118     const uint8x16x2_t pairs[4] = {
119         { .val = { m[0].val[0], m[1].val[0] } },
120         { .val = { m[0].val[1], m[1].val[1] } },
121         { .val = { m[0].val[2], m[1].val[2] } },
122         { .val = { m[0].val[3], m[1].val[3] } },
123     };
124
125    m[0].val[0] = vqtbl2q_u8(pairs[0], pack_shf_inv.val[0]);
126    m[0].val[1] = vqtbl2q_u8(pairs[1], pack_shf_inv.val[0]);
127    m[0].val[2] = vqtbl2q_u8(pairs[2], pack_shf_inv.val[0]);
128    m[0].val[3] = vqtbl2q_u8(pairs[3], pack_shf_inv.val[0]);
129
130    m[1].val[0] = vqtbl2q_u8(pairs[0], pack_shf_inv.val[1]);

```

```

131     m[1].val[1] = vqtbl2q_u8(pairs[1], pack_shf_inv.val[1]);
132     m[1].val[2] = vqtbl2q_u8(pairs[2], pack_shf_inv.val[1]);
133     m[1].val[3] = vqtbl2q_u8(pairs[3], pack_shf_inv.val[1]);
134
135     // take care not to shift mask bits out of the register
136     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[1], pack_mask_0, 1);
137     gift_64_vec_sliced_swapmove(&m[0].val[2], &m[0].val[3], pack_mask_0, 1);
138     gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[1], pack_mask_0, 1);
139     gift_64_vec_sliced_swapmove(&m[1].val[2], &m[1].val[3], pack_mask_0, 1);
140
141     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[0].val[2], pack_mask_1, 2);
142     gift_64_vec_sliced_swapmove(&m[0].val[1], &m[0].val[3], pack_mask_1, 2);
143     gift_64_vec_sliced_swapmove(&m[1].val[0], &m[1].val[2], pack_mask_1, 2);
144     gift_64_vec_sliced_swapmove(&m[1].val[1], &m[1].val[3], pack_mask_1, 2);
145
146     // make bytes (a0 b0 c0 d0 a4 b4 c4 d4 -> a0 b0 c0 d0 e0 f0 g0 h0)
147     gift_64_vec_sliced_swapmove(&m[0].val[0], &m[1].val[0], pack_mask_2, 4);
148     gift_64_vec_sliced_swapmove(&m[0].val[2], &m[1].val[2], pack_mask_2, 4);
149     gift_64_vec_sliced_swapmove(&m[0].val[1], &m[1].val[1], pack_mask_2, 4);
150     gift_64_vec_sliced_swapmove(&m[0].val[3], &m[1].val[3], pack_mask_2, 4);
151 }
152
153 void gift_64_vec_sliced_subcells(uint8x16x4_t cs[restrict 2])
154 {
155     cs[0].val[1] = veorq_u8(cs[0].val[1],
156                             vandq_u8(cs[0].val[0], cs[0].val[2]));
157     uint8x16_t t = veorq_u8(cs[0].val[0],
158                             vandq_u8(cs[0].val[1], cs[0].val[3]));
159     cs[0].val[2] = veorq_u8(cs[0].val[2], vorrq_u8(t, cs[0].val[1]));
160     cs[0].val[0] = veorq_u8(cs[0].val[3], cs[0].val[2]);
161     cs[0].val[1] = veorq_u8(cs[0].val[1], cs[0].val[0]);
162     cs[0].val[0] = vmvnq_u8(cs[0].val[0]);
163     cs[0].val[2] = veorq_u8(cs[0].val[2], vandq_u8(t, cs[0].val[1]));
164     cs[0].val[3] = t;
165
166     cs[1].val[1] = veorq_u8(cs[1].val[1],
167                             vandq_u8(cs[1].val[0], cs[1].val[2]));
168     t = veorq_u8(cs[1].val[0],
169                  vandq_u8(cs[1].val[1], cs[1].val[3]));
170     cs[1].val[2] = veorq_u8(cs[1].val[2], vorrq_u8(t, cs[1].val[1]));
171     cs[1].val[0] = veorq_u8(cs[1].val[3], cs[1].val[2]);
172     cs[1].val[1] = veorq_u8(cs[1].val[1], cs[1].val[0]);
173     cs[1].val[0] = vmvnq_u8(cs[1].val[0]);
174     cs[1].val[2] = veorq_u8(cs[1].val[2], vandq_u8(t, cs[1].val[1]));
175     cs[1].val[3] = t;
176 }
177
178 void gift_64_vec_sliced_subcells_inv(uint8x16x4_t cs[restrict 2])
179 {
180     uint8x16_t t = cs[0].val[3];
181     cs[0].val[2] = veorq_u8(cs[0].val[2], vandq_u8(t, cs[0].val[1]));
182     cs[0].val[0] = vmvnq_u8(cs[0].val[0]);
183     cs[0].val[1] = veorq_u8(cs[0].val[1], cs[0].val[0]);
184     cs[0].val[3] = veorq_u8(cs[0].val[0], cs[0].val[2]);
185     cs[0].val[2] = veorq_u8(cs[0].val[2], vorrq_u8(t, cs[0].val[1]));
186     cs[0].val[0] = veorq_u8(t, vandq_u8(cs[0].val[1], cs[0].val[3]));
187     cs[0].val[1] = veorq_u8(cs[0].val[1],
188                             vandq_u8(cs[0].val[0], cs[0].val[2]));
189
190     t = cs[1].val[3];
191     cs[1].val[2] = veorq_u8(cs[1].val[2], vandq_u8(t, cs[1].val[1]));
192     cs[1].val[0] = vmvnq_u8(cs[1].val[0]);

```

```

193     cs[1].val[1] = veorq_u8(cs[1].val[1], cs[1].val[0]);
194     cs[1].val[3] = veorq_u8(cs[1].val[0], cs[1].val[2]);
195     cs[1].val[2] = veorq_u8(cs[1].val[2], vorrq_u8(t, cs[1].val[1]));
196     cs[1].val[0] = veorq_u8(t, vandq_u8(cs[1].val[1], cs[1].val[3]));
197     cs[1].val[1] = veorq_u8(cs[1].val[1],
198                             vandq_u8(cs[1].val[0], cs[1].val[2]));
199 }
200
201 void gift_64_vec_sliced_permute(uint8x16x4_t cs[restrict 2])
202 {
203     cs[0].val[0] = vqtblq_u8(cs[0].val[0], perm.val[0]);
204     cs[0].val[1] = vqtblq_u8(cs[0].val[1], perm.val[1]);
205     cs[0].val[2] = vqtblq_u8(cs[0].val[2], perm.val[2]);
206     cs[0].val[3] = vqtblq_u8(cs[0].val[3], perm.val[3]);
207
208     cs[1].val[0] = vqtblq_u8(cs[1].val[0], perm.val[0]);
209     cs[1].val[1] = vqtblq_u8(cs[1].val[1], perm.val[1]);
210     cs[1].val[2] = vqtblq_u8(cs[1].val[2], perm.val[2]);
211     cs[1].val[3] = vqtblq_u8(cs[1].val[3], perm.val[3]);
212 }
213
214 void gift_64_vec_sliced_permute_inv(uint8x16x4_t cs[restrict 2])
215 {
216     cs[0].val[0] = vqtblq_u8(cs[0].val[0], perm_inv.val[0]);
217     cs[0].val[1] = vqtblq_u8(cs[0].val[1], perm_inv.val[1]);
218     cs[0].val[2] = vqtblq_u8(cs[0].val[2], perm_inv.val[2]);
219     cs[0].val[3] = vqtblq_u8(cs[0].val[3], perm_inv.val[3]);
220
221     cs[1].val[0] = vqtblq_u8(cs[1].val[0], perm_inv.val[0]);
222     cs[1].val[1] = vqtblq_u8(cs[1].val[1], perm_inv.val[1]);
223     cs[1].val[2] = vqtblq_u8(cs[1].val[2], perm_inv.val[2]);
224     cs[1].val[3] = vqtblq_u8(cs[1].val[3], perm_inv.val[3]);
225 }
226
227 void gift_64_vec_sliced_generate_round_keys(uint8x16x4_t rks[restrict ROUNDS_GIFT_64
228 ] [2],
229                                             const uint64_t key[restrict 2])
230 {
231     uint64_t key_state[] = {key[0], key[1]};
232     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
233         const int v = (key_state[0] >> 0) & 0xffff;
234         const int u = (key_state[0] >> 16) & 0xffff;
235
236         // add round key (RK=U||V)
237         // (slice 2 stays unused)
238         uint64_t rk[6] = { 0x0UL };
239         for (size_t i = 0; i < 8; i++) {
240             int key_bit_v = (v >> (i + 0)) & 0x1;
241             int key_bit_u = (u >> (i + 0)) & 0x1;
242             rk[0] ^= (uint64_t)key_bit_v << (i * 8);
243             rk[2] ^= (uint64_t)key_bit_u << (i * 8);
244
245             key_bit_v = (v >> (i + 8)) & 0x1;
246             key_bit_u = (u >> (i + 8)) & 0x1;
247             rk[1] ^= (uint64_t)key_bit_v << (i * 8);
248             rk[3] ^= (uint64_t)key_bit_u << (i * 8);
249         }
250
251         // add single bit
252         rk[5] ^= 1UL << (7 * 8);
253
254         // add round constants

```

```

254         rk[4] ^= ((uint64_t)(round_const[round] >> 0) & 0x1) << (0 * 8);
255         rk[4] ^= ((uint64_t)(round_const[round] >> 1) & 0x1) << (1 * 8);
256         rk[4] ^= ((uint64_t)(round_const[round] >> 2) & 0x1) << (2 * 8);
257         rk[4] ^= ((uint64_t)(round_const[round] >> 3) & 0x1) << (3 * 8);
258         rk[4] ^= ((uint64_t)(round_const[round] >> 4) & 0x1) << (4 * 8);
259         rk[4] ^= ((uint64_t)(round_const[round] >> 5) & 0x1) << (5 * 8);
260
261         // extend bits to bytes
262         for (size_t i = 0; i < 6; i++) {
263             rk[i] |= rk[i] << 1;
264             rk[i] |= rk[i] << 2;
265             rk[i] |= rk[i] << 4;
266         }
267
268         rks[round][0].val[0] = vsetq_lane_u64(rk[0], rks[round][0].val[0], 0);
269         rks[round][0].val[0] = vsetq_lane_u64(rk[1], rks[round][0].val[0], 1);
270         rks[round][0].val[1] = vsetq_lane_u64(rk[2], rks[round][0].val[1], 0);
271         rks[round][0].val[1] = vsetq_lane_u64(rk[3], rks[round][0].val[1], 1);
272         rks[round][0].val[2] = vdupq_n_u8(0);
273         rks[round][0].val[3] = vsetq_lane_u64(rk[4], rks[round][0].val[3], 0);
274         rks[round][0].val[3] = vsetq_lane_u64(rk[5], rks[round][0].val[3], 1);
275         rks[round][1] = rks[round][0];
276
277         // update key state
278         int k0 = (key_state[0] >> 0) & 0xffffUL;
279         int k1 = (key_state[0] >> 16) & 0xffffUL;
280         k0 = (k0 >> 12) | ((k0 & 0xfff) << 4);
281         k1 = (k1 >> 2) | ((k1 & 0x3) << 14);
282         key_state[0] >>= 32;
283         key_state[0] |= (key_state[1] & 0xffffffffUL) << 32;
284         key_state[1] >>= 32;
285         key_state[1] |= ((uint64_t)k0 << 32) | ((uint64_t)k1 << 48);
286     }
287 }
288
289 void gift_64_vec_sliced_init(void)
290 {
291     // bit packing shuffle
292     pack_shf = vld1q_u8_x2((uint8_t*)&pack_shf_u64[0]);
293
294     // inverse bit packing shuffle
295     pack_shf_inv = vld1q_u8_x2((uint8_t*)&pack_shf_inv_u64[0]);
296
297     // permutations
298     perm = vld1q_u8_x4((uint8_t*)&perm_u64[0]);
299
300     // inverse permutations
301     perm_inv = vld1q_u8_x4((uint8_t*)&perm_inv_u64[0]);
302
303     // packing masks
304     pack_mask_0 = vdupq_n_u8(0x55);
305     pack_mask_1 = vdupq_n_u8(0x33);
306     pack_mask_2 = vdupq_n_u8(0x0f);
307 }
308
309 void gift_64_vec_sliced_encrypt(uint64_t c[restrict 16],
310                                const uint64_t m[restrict 16],
311                                const uint8x16x4_t rks[restrict ROUNDS_GIFT_64][2])
312 {
313     uint8x16x4_t s[2];
314     s[0] = vld1q_u8_x4((uint8_t*)&m[0]);
315     s[1] = vld1q_u8_x4((uint8_t*)&m[8]);

```

```

316     gift_64_vec_sliced_bits_pack(s);
317
318     for (int round = 0; round < ROUNDS_GIFT_64; round++) {
319         gift_64_vec_sliced_subcells(s);
320         gift_64_vec_sliced_permute(s);
321
322         // round key addition
323         s[0].val[0] = veorq_u8(s[0].val[0], rks[round][0].val[0]);
324         s[0].val[1] = veorq_u8(s[0].val[1], rks[round][0].val[1]);
325         s[0].val[2] = veorq_u8(s[0].val[2], rks[round][0].val[2]);
326         s[0].val[3] = veorq_u8(s[0].val[3], rks[round][0].val[3]);
327         s[1].val[0] = veorq_u8(s[1].val[0], rks[round][1].val[0]);
328         s[1].val[1] = veorq_u8(s[1].val[1], rks[round][1].val[1]);
329         s[1].val[2] = veorq_u8(s[1].val[2], rks[round][1].val[2]);
330         s[1].val[3] = veorq_u8(s[1].val[3], rks[round][1].val[3]);
331     }
332
333     gift_64_vec_sliced_bits_unpack(s);
334     vst1q_u8_x4((uint8_t*)&c[0], s[0]);
335     vst1q_u8_x4((uint8_t*)&c[8], s[1]);
336 }
337
338 void gift_64_vec_sliced_decrypt(uint64_t m[restrict 16],
339                               const uint64_t c[restrict 16],
340                               const uint8x16x4_t rks[restrict ROUNDS_GIFT_64][2])
341 {
342     uint8x16x4_t s[2];
343     s[0] = vld1q_u8_x4((uint8_t*)&c[0]);
344     s[1] = vld1q_u8_x4((uint8_t*)&c[8]);
345     gift_64_vec_sliced_bits_pack(s);
346
347     for (int round = ROUNDS_GIFT_64 - 1; round >= 0; round--) {
348         // round key addition
349         s[0].val[0] = veorq_u8(s[0].val[0], rks[round][0].val[0]);
350         s[0].val[1] = veorq_u8(s[0].val[1], rks[round][0].val[1]);
351         s[0].val[2] = veorq_u8(s[0].val[2], rks[round][0].val[2]);
352         s[0].val[3] = veorq_u8(s[0].val[3], rks[round][0].val[3]);
353         s[1].val[0] = veorq_u8(s[1].val[0], rks[round][1].val[0]);
354         s[1].val[1] = veorq_u8(s[1].val[1], rks[round][1].val[1]);
355         s[1].val[2] = veorq_u8(s[1].val[2], rks[round][1].val[2]);
356         s[1].val[3] = veorq_u8(s[1].val[3], rks[round][1].val[3]);
357
358         gift_64_vec_sliced_permute_inv(s);
359         gift_64_vec_sliced_subcells_inv(s);
360     }
361
362     gift_64_vec_sliced_bits_unpack(s);
363     vst1q_u8_x4((uint8_t*)&m[0], s[0]);
364     vst1q_u8_x4((uint8_t*)&m[8], s[1]);
365 }

```

Appendix B

Lorem dolor

Bibliography

- [1] Subhadeep Banik et al. “GIFT: A Small Present”. In: Aug. 2017, pp. 321–345. ISBN: 978-3-319-66786-7. DOI: [10.1007/978-3-319-66787-4_16](https://doi.org/10.1007/978-3-319-66787-4_16).
- [2] Kazumaro Aoki et al. “Camellia: A 128-Bit Block Cipher Suitable for Multiple Platforms — Design and Analysis”. In: *Selected Areas in Cryptography*. Ed. by Douglas R. Stinson and Stafford Tavares. Berlin, Heidelberg: Springer Berlin Heidelberg, 2001, pp. 39–56. ISBN: 978-3-540-44983-6.
- [3] ARM Limited. “The Future is Built on Arm”. URL: <https://www.arm.com/company> (visited on 04/07/2023).
- [4] “ODROID-N2+ with 4GByte RAM”. URL: <https://www.hardkernel.com/shop/odroid-n2-with-4gbyte-ram-2/#tab-description> (visited on 04/03/2023).
- [5] ARM Limited. “ARM Architecture Reference Manual. ARMv8, for ARMv8-A architecture profile”. 2013. URL: <https://developer.arm.com/documentation/ddi0487/latest/> (visited on 02/02/2023).
- [6] ARM Limited. “Arm Neon Intrinsics Reference”. 2022. URL: https://arm-software.github.io/acle/neon_intrinsics/advsimd.html (visited on 02/07/2023).
- [7] Ryad Benadjila et al. “Implementing Lightweight Block Ciphers on x86 Architectures”. In: *Selected Areas in Cryptography – SAC 2013*. Ed. by Tanja Lange, Kristin Lauter, and Petr Lisoněk. Berlin, Heidelberg: Springer Berlin Heidelberg, 2014, pp. 324–351. ISBN: 978-3-662-43414-7.
- [8] ARM Limited. “Cortex-A72 Software Optimization Guide”. 2015. URL: <https://developer.arm.com/documentation/uan0016/a/> (visited on 04/06/2023).